

派業歷

# **PyElly User's Manual**

Release v1.0

14 December 2014

Clinton P. Mah

Walnut Creek, CA 94595

## Table of Contents

1. Introduction	5
2. The Syntax of a Language	9
3. The Semantics of a Language	13
4. Defining Tables of PyElly Rules	17
4.1 Grammar (A.g.elly)	19
4.1.1 Syntactic Rules	19
4.1.2 Grammar-Defined Words	20
4.1.3 Generative Semantic Subprocedures	21
4.1.4 Global Variable Initializations	21
4.2 Special Patterns (A.p.elly)	22
4.3 Macro Substitutions (A.m.elly)	25
5. Operations for PyElly Generative Semantics	27
5.1 Insertion of Strings	27
5.2 Subroutine Linkage	27
5.3 Buffer Management	28
5.4 Local Variable Operations	28
5.5 Local Variable Set Operations	30
5.6 Global Variable Operations	30
5.7 Control Structures	30
5.8 Character Manipulation	32
5.9 Selection From a Table	34
5.10 Buffer Searching	35
5.11 Execution Monitoring	35
5.12 Capitalization	36
5.13 Semantic Subprocedure Invocation	36
6. Simple PyElly Programming Examples	37
6.1 Default Semantic Procedures	37
6.2 A Simple Grammar Example	38
7. Running PyElly From a Command Line	42
8. Advanced Capabilities: Grammar	46
8.1 Syntactic Features	46
8.2 The ... Syntactic Type	48

## Table of Contents

9. Advanced Capabilities: Vocabulary	51
9.1 Working with the UNKN Syntactic Type	51
9.2 Breaking Down Unknown Words	52
9.2.1 Inflectional Stemming	52
9.2.2 Morphology	55
9.2.2.1 Word Endings	55
9.2.2.2 Word Beginnings	57
9.3 Entity Extraction	58
9.3.1 Numbers	59
9.3.2 Dates and Times	59
9.3.3 Other Entities	60
9.4 PyElly Vocabulary Tables	61
10. Logic for PyElly Cognitive Semantics	65
10.1 The Form of Cognitive Semantic Clauses	67
10.2 Cognitive Semantic Elements	68
10.2.1 Fixed Scoring	68
10.2.2 Semantic Features	69
10.2.2.1 Semantic Features in Cognitive Semantic Clauses	69
10.2.2.2 Semantic Features in Generative Semantics	70
10.2.3 Semantic Concepts	70
10.2.3.1 Concepts in Cognitive Semantic Logic	71
10.2.3.2 Semantic Concepts in Language Definition Files	73
10.2.3.2.1 Conceptual Hierarchy Definition (A.h.elly)	73
10.2.3.2.2 Semantic Concepts in Grammar Rules (A.g.elly)	74
10.2.3.2.3 Semantic Concepts in Vocabulary Table Entries (A.v.elly)	75
11. Sentences and Punctuation	76
11.1 Extending Sentence Recognition	77
11.1.1 Stop Punctuation Exceptions	77
11.1.2 Exotic Punctuation	78
11.2 Predefined Punctuation Syntax	78
12. PyElly Parsing	80
12.1 A Bottom-Up Framework	80

## Table of Contents

12.2 Token Lookup	81
12.3 Building a Parse Tree	83
12.3.1 Context-Free Analysis Main Loop	83
12.3.2 Special Modifications	84
12.3.3 Type 0 Extensions	84
12.4 Success and Failure in Parsing	85
12.5 Parse Data Dumps and Tree Diagrams	85
12.6 Resource Limits	89
13. Developing Language Rules and Troubleshooting	90
13.1 Pre-Checks on Language Rule Files	90
13.2 A General Application Development Approach	92
13.3 Miscellaneous Tips	92
14. PyElly Applications	97
Appendix A. Python Implementation	102
Appendix B. Historical Background	107
Appendix C. Berkeley Database	109
Appendix D. PyElly System Testing	111

# 1. Introduction

PyElly is an open-source software tool for creating computer scripts to analyze and rewrite English and other natural language text. This processing will of course fall far short of realizing the talking robot fantasies of Hollywood, but with only modest effort, you can still build many nontrivial linguistic applications short of full understanding and take care of low-level language details that can often get in the way of solving more interesting problems involving text data.

PyElly can also be helpful if you just want to gain experience with the nitty-gritty of language processing. You can quickly build scripts to do basic tasks like conjugating French verbs, rephrasing information requests into a formal query language, compressing messages for texting, extracting names and other entities from a text stream, and even re-creating the storied Doctor simulation of Rogerian psychoanalysis.

We have been building such natural language applications since our machines were a million times less powerful than they are now. The overall problem here remains quite challenging, however, and even today's technology may require you to work hard to develop algorithms and codify linguistic knowledge just to do something quite simple. PyElly aims to make that work easier through ready-made tools and resources, all in a single free open-source package.

Why do we need yet another natural language processing toolkit? To begin with, a complete natural language solution is still far off, and so we can benefit from having a diversity of tested methods pertinent to the problem. Also, though PyElly is all new code, it is really a legacy system, with its core components dating back about 40 years. This sounds quite ancient, but language changes slowly, and mature software tools can make it easier to work with text data.

The impetus for PyElly and its predecessors came from observing that many natural language systems tend to run into the same everyday subproblems. For example, information retrieval and machine learning with text data can work better when we can reduce its words of the text into their roots. Instead of contending with variants like `RELATION`, `RELATIONAL`, `RELATIVELY`, and `RELATING`, we have just `RELATE`. This is of course the familiar stemming problem, but available free resources to do such associations have often been disappointing.

A stemmer of course is not hard to build, but it takes time and commitment to do a good job, and one hardly wants to do this from scratch for every new project. That is true of other basic language processing capabilities as well. So, it seems logical to pull together at least some kind of reusable software library here, but it would be even more helpful if we could integrate our tools and resources more closely. PyElly does that.

The current implementation of PyElly is intended primarily for educational use and so was written entirely in Python, currently a favored first programming language in high schools. This should allow students to adapt and incorporate PyElly into class projects that have to be completed fairly quickly. PyElly can be of broader interest, though,

because of its range of natural language support: stemming, tokenizing, entity recognition, sentence extraction, idiomatic transformation, rule-driven syntactic analysis, and semantic ambiguity handling.

Effective use of PyElly will of course still require some linguistic expertise. You will have to be able to define the details of the language processing that you want, but much of the basics here have been prebuilt in PyElly if you are working with English input. The standard PyElly distribution includes the definition scripts for six different actual example applications that you can modify as a head start in constructing your own.

The current PyElly package consists of a series of Python modules in fifty-eight source files. The code should run on any computer platform with a Python 2.7 interpreter, including Windows 7 and 8, Linux, Mac OS X and other flavors of Unix, IOS for iPhone and iPad, and Android. The PyElly source is downloadable from GitHub under a standard BSD license; you may freely modify and extend it as needed. Though intended mainly for education, there are no restrictions on commercial usage.

For recognizing just a few dozen sentences, PyElly is probably overkill; you could handle these directly by writing custom code in any standard programming language. More often, however, possible input sentences are too many to list out fully, and you will have to characterize them more generally through various rules describing how the words you expect to see are formed, how they combine in text, and how they are to be interpreted.

PyElly is set up as a kind of translator: it reads in, analyzes, and writes out transformed text according to the rules that you supply. So an English sentence like “She goes slowly” might be rewritten in French as “Elle va lentement” or in traditional Chinese 她走慢慢. Or you might reduce the original sentence to just “slow” by stripping out suffixes and words of low content. Or you may want to rephrase the sentence as a question like “Does she go slowly?” All of this can be done entirely within PyElly.

PyElly rules will be of various types. The main ones will define a grammar and vocabulary for the sentences of an input language plus associated semantic procedures for rewriting those sentences to get a desired output. Creating such rules requires some trial and error, but usually should be no more difficult than setting up macros in a word processor. PyElly will get you started quickly here and also provide debugging aids to help track down any problems.

Many natural language system building tools, especially those in academic research, tend to address the most thorny problems in language interpretation. These are opportunities for impressive processing gymnastics and often lead to interesting theoretical papers without necessarily producing practical results. PyElly tries here to be simple and pragmatic instead. In response to classically tough sentences like “Time flies like an arrow,” it is all right just to respond with “Huh?”

PyElly was built to be compact enough to run on mobile devices, if necessary. Excluding the Python environment, compiled PyElly code along with encoded rules and other data for an application should typically require less than 500 Kbytes of storage, depending on the number of rules actually defined. A major project may involve hundreds of grammar

rules and thousands of vocabulary elements, but some useful text analyses require just a few dozen rules and no explicit domain vocabulary.

What is a grammar, and what is a vocabulary? A vocabulary establishes the range of words you want to recognize; a grammar defines how those words can be arranged into sentences of interest. You may also specify idiomatic rewriting of particular input word sequences prior to analysis as well as define patterns to make sense of various classes of unknown words. For example, you can recognize 800 telephone numbers or Russian surnames ending in -OV without having to list them all out.

This manual will explain how to do all of this and also cover some basics of language and language processing that every PyElly user should know. To use PyElly, you should already be able to create and edit text files from the command line of whatever system you choose to work on and set up file directories where PyElly can find them. In an ideal world, an interactive development environment (IDE) could make everything easier here, but that is yet to be.

Currently, PyElly is biased toward English input, although it can read the entire Latin-1 subset of Unicode as input. That subset includes the familiar ASCII characters as well as all the letters with diacritical marks used in Western European languages. For example, PyElly knows that é is a vowel, that ß is a letter, and that Æ is the uppercase form of œ. This can be helpful even for nominally English data, since we often encounter terms with foreign spellings like NÉE.

As any beginning student of a foreign language soon learns, the rules are often messy. Irregularities always arise to trip up someone mechanically trying to speak or write from a simple grammar. PyElly users will face the same kind of problem here, but by working generally at first and dealing with exceptions as they crop up, we can evolve our rules to reach eventually some useful level of parlance. There is no royal road to natural language understanding, but persistence can achieve wonders over time, and PyElly can help to keep you going in such a long-term endeavor.

You need not be an experienced linguist or computer programmer to develop PyElly applications; and I have tried to write this manual to be understandable by non-experts. The only requirements for users are basic computer literacy as expected of 21st-Century high school graduates, linguistic knowledge as might be picked up from a first course in a foreign language, and willingness to learn. You should start out with simple kinds of PyElly processing and progress to more complex analysis with more experience.

In addition to this introduction, the PyElly User's Manual consists of thirteen other major sections plus four appendices. Sections 2 through 7 should be read in sequence as they are a tutorial on how get started quickly on using PyElly. Sections 8 through 11 deal with advanced features that may be helpful for specific kinds of applications. Section 12 explains PyElly parsing, Section 13 lays out some strategies for application development, and Section 14 describes some practical PyElly applications.

PyElly ("Python Elly") was inspired by the Eliza system created by Joseph Weizenbaum over 50 years ago for studying natural language conversation, but PyElly has a

completely different genesis. Its Python implementation is the latest in a series of related natural language processors going back four generations: Jelly (Java, 1999), nlf (C, 1984 and 1994), the Adaptive Query Facility (FORTRAN, 1981), and PARLEZ (PDP-11 assembly language, 1977). The PyElly parsing algorithm and the ideas of cognitive and generative semantics are based on Vaughn Pratt's LINGOL (LISP, 1973). Frederick Thompson's REL system (1972) also influenced PyElly.

The PyElly website is at <https://sites.google.com/site/pyellynaturallanguage/> .



## 2. The Syntax of a Language

A language is as a way of putting together words or other symbols to form sentences that other people can make sense of in some context. In general, not all combinations of symbols will make a meaningful sentence; for example, “Cat the when” is nonsense in English. To define a language that you want PyElly to process, you must first identify those combinations of symbols that do make sense and then assign suitable interpretations to them.

If a language is small enough, such as the repertory of obscene gestures, we can simply list all its possible “sentences” and write down what each of them mean. Most nontrivial languages, though, have so many possible sentences that this approach is impractical. Instead one must note that languages tend to have regular structures; and by identifying those structures, we can formally characterize the language much more concisely than by listing all possible sentences one after another.

The structural description of a language is called a grammar. It states what the building blocks of a language are and how they form simple structures, which in turn combine into successively more complex structures. Almost everyone has studied grammar in school, but formal grammars go into much more detail. They commonly are organized as sets of syntactic rules describing particular kinds of language structure in sentences. Such syntactic rules will provide a basis for both generating and recognizing sentences of a language with a computer.

In linguistics, a formal rule of grammar is expressed in terms of how one or more structures come together to produce a new composite structure. These syntactic rules are commonly written with an arrow notation as follows:

$$W \rightarrow X \ Y \ Z$$

This rule states that a *W*-structure can be composed of an *X*-structure followed by a *Y*-structure followed by a *Z*-structure; for example, a noun phrase can consist of a number, followed by an adjective, and followed by a noun:

$$\text{NOUNPHRASE} \rightarrow \text{NUMBER} \ \text{ADJECTIVE} \ \text{NOUN}$$

There is nothing mysterious here; it is like the kind of sentence diagramming once taught in junior high school and now coming back into vogue. In fact, we could draw the following equivalent diagram on a blackboard for the syntactic rule above.



where the `W` can in turn be part of a higher-level structure and `X`, `Y`, and `Z` can split out into further substructures. Either way of describing a syntactic structure is fine, but the arrow notation will be more compact and easier to type out on a keyboard, especially as syntax grows more complex.

Syntactic rules generally can be much more complicated than `W->X Y Z`, but with PyElly, it turns out that we can get by with rules restricted to just three types:

```
X->word
X->Y
X->Y Z
```

where `X`, `Y`, and `Z` are structure types and `word` is a word or other kind of vocabulary element.

For example, to express a more complex rule like

```
R->A B C
```

we can instead use the pair of restricted rules

```
R->A T
T->B C
```

where `T` is a unique intermediate structure introduced here solely to stand for a `B` followed by a `C`.

Here is a set of restricted grammar rules that might be employed to describe the structure of the sentence “It is red.”

```
SENTENCE->SUBJECT PREDICATE
SUBJECT->PRONOUN
PRONOUN->it
PREDICATE->COPULA ADJECTIVE
COPULA->is
ADJECTIVE->red
```

This uses all three types of PyElly restricted grammar rules.

The structure of a sentence as implied by these rules can be expressed graphically as a labeled tree diagram, where the root type must be `SENTENCE` and where branching corresponds to splitting into constituent substructures. By convention, the tree is always upside down, and the bottom of the tree will show the individual words of the sentence. For example, the sentence “It is red” would have the following diagram here:



The derivation of such a diagram for a sentence from a given set of syntactic rules is called “parsing.” The diagram itself is called a “parse tree,” and its labeled parts are called the “phrase nodes” of the parse tree; for example, the phrase node `PREDICATE` in the tree above encompasses the actual sentence phrase “is red.” Given the grammar rules above, PyElly can carry out this analysis automatically for our sentence; and the resulting tree diagram then provides a starting point for interpreting the sentence.

Our grammar above describes only a single sentence, but we can extend the range of its coverage by adding more rules for new kinds of structures and vocabulary. For example, the new rules

```

SUBJECT->DETERMINER NOUN
DETERMINER->an
NOUN->apple

```

will describe “An apple is red” as a sentence for PyElly to process. We can continue to build up our grammar here. The rules

```

PREDICATE->VERB
VERB->falls

```

will also put “It falls” and “An apple falls” into the language we can recognize. This expansion of rules may be continued to encompass still more types of structure and still more vocabulary. You can add such new rules in any order you want.

The key idea here is that a limited number of rules can be combined in various ways to describe many different sentences. There is still the problem of choosing the proper mix of rules here to describe a language in the most natural and efficient way, but we do fairly well by simply adding one or two rules at a time as done above. In sophisticated applications, we may eventually need hundreds of such rules, but these can still be worked out in small steps.

Technically speaking, PyElly grammar rules as described here define a “context-free language.” With such a grammar, you cannot correlate the possibilities for a given structure with the possibilities for a parallel structure in a parse tree. For example, consider the context-free rule with the parallel structures `SUBJECT` and `PREDICATE`.

```
SENTENCE->SUBJECT PREDICATE
```

In languages like English, subjects and predicates have to agree with each other according to the attributes of person and number: “We fall” versus “he falls”. When grammatically acceptable `SUBJECT` and `PREDICATE` structures can be formed in more than one way, our context-free rule here by itself will not allow us to restrict a `SENTENCE` to have only certain combinations of subjects and predicates for agreement. We could of course write explicitly correlated rules like

```
SENTENCE->SUBJECT1 PREDICATE1
SENTENCE->SUBJECT2 PREDICATE2
SENTENCE->SUBJECT3 PREDICATE3
```

where `SUBJECTi` would always agree properly with `PREDICATEi`, but this has the disadvantage of greatly multiplying the number of rules we have to define. A good natural language toolkit should make our job easier than this.

Although English and other natural languages are not context-free in the theoretical sense, we still want to treat them that way from a practical perspective. The advantage in doing so is that we can then apply many sophisticated techniques developed for parsing artificial computer languages, which do tend to be context-free. This is in fact the approach implemented in PyElly and its predecessors.

For convenience, PyElly also incorporates semantic checking of the results of parsing and allows some shortcuts to make grammars more concise (see Section 8). These extensions can be put on top of a context-free parser to give it some context-sensitive capabilities, although some kinds of sentences still cannot be handled by PyElly. (The classic context-sensitive examples are parallel subjects and predicates, such as in the sentence “He and she got cologne and perfume, respectively.”)

The syntax of natural language can get quite complex in general; but we usually can break this down into simpler structures. The challenge of defining a PyElly grammar is to capture enough of such simpler structures in grammar rules to support a proper analysis of any input sentence that we are likely to see.

You must be able to understand most of the discussion in this section in order to proceed further with PyElly. A good text for those interested in learning more about language and formal grammars is John Lyon's book *Introduction to Theoretical Linguistics* (Cambridge University Press, 1968). This is written for college-level readers, but sticks to the basics that you will need to know.

### 3. The Semantics of a Language

The notion of meaning has always been difficult to talk about. It can be complicated even for individual sentences in a language, because meaning involves not only their grammatical structure, but also where it is used and who is using it. A simple expression like “Thank you” can take on different significance, depending on whether the speaker is a thug collecting extortion money, the senior correspondent at a White House news conference, or a disaster victim after an arduous rescue.

Practical computer natural language applications cannot deal with all the potential meanings of sentences, since this would require modeling almost everything in a person's view of world and self. A more realistic approach is to ask what meanings will actually be appropriate for a computer to understand in a particular application. If the role of a system in a user organization is to provide, say, only information about employee benefits from a policy manual, then it probably has no reason to handle references to subjects like sex, golf, or the current weather.

Here we shall limit the scope of semantics even further: PyElly will deal with the meaning of sentences only to the extent of being able to translate them into some another language and to evaluate alternate options when we have more than one possible translation. This has the advantage of making semantics less mysterious while allowing us still to achieve useful kinds of language processing.

For example, the meaning of the English sentence “I love you” could be expressed in French as “Je t'aime.” Or we might translate the English “How much does John earn?” into a data base query language “SELECT SALARY FROM PAYROLL WHERE EMPLOYEE=JOHN.” Or we could convert the statement “I feel hot” into a computer command line like

```
set thermostat /relative /fahrenheit -5
```

In a sense, we have cheated here, avoiding the problem of meaning in one language by passing it off to another language. Such a translation, however, can be quite useful if we already happen to have a processor that understands the second language, but not the first. This is definitely a modest approach to semantics; but it beats talking endlessly about the philosophical meaning of meaning without ever accomplishing anything.

As noted before, the large number of possible sentences in a natural language prevents us from compiling a table to map each input into its corresponding output. Instead, we must break the problem down and take the semantics of the various constituent structures defined by a grammar and combine their individual interpretations to derive the overall meaning of a given sentence.

With PyElly, we define the semantics of a language structure as procedures associated with the grammatical rule for the structure. There will actually be two different kinds of semantic procedures here: those for writing out translations will be called “generative,” while those for evaluating alternative translations will be called “cognitive.” At this

stage, however, we shall focus on the generative part, and leave cognitive part to Section 10 since they operate quite differently.

A successful PyElly sentence analysis will produce a parse tree describing its syntactic structure. Each phrase node of that tree will be due to a particular grammatical rule, and associated with that rule will be a generative semantic procedure defining its meaning. You will have to supply such a procedure for each of your grammar rules, though often you can just take the defaults defined by PyElly for its three types of grammar rules.

The top phrase node of a parse tree should always be that of the type `SENTENCE`. The generative procedure for that phrase node will be called by PyElly to begin the overall translation of the original input sentence. This should then set off a cascade of other procedure calls through the various lower constituent structures of the sentence to produce a final output. The actual ordering of calls to subconstituent procedures will be determined by the logic of the procedures at each level of the tree.

Each generative semantic procedure basically will work on the text characters in a series of output buffers. This will involve standard text editing operations commonly supported in word processing programs: inserting and deleting, buffer management, searching, substitution, and transfers between buffers. Consistent with PyElly semantics being procedures, there will also be local and global variables, structured programming control structures, subprocedures, simple lookup, and set manipulation.

Communication between different semantic procedures will be through local and global variables. The value of any such variable will always be a string of arbitrary Unicode characters, possibly the null string. Global variables will be accessible to all procedures and will remain defined even across the processing of successive sentences, serving as a long-term memory for PyElly translations.

Local variables will have a limited scope as in programming languages like C or PASCAL. They are defined in the procedure where they are declared and also in those procedures called as subroutines either directly or indirectly. When there are multiple active declarations of a variable with a given name visible to a semantic procedure, the most recent one applies. Upon exit from a procedure, all of its local variables immediately become undefined.

Here are some semantic procedures to illustrate how a PyElly translation might actually work. Suppose that we define the five grammar rules

```
SENTENCE->SUBJECT PREDICATE
SUBJECT->PRONOUN
PRONOUN->we
PREDICATE->VERB
VERB->know
```

With these rules, we can implement a simple translator from English into French with the five semantic procedures below, defined respectively for each rule above. For the

time being, the commands in the procedures will be expressed in ordinary English. Assume for now that such commands will control the entry of text into some output area, such as a text field in a window of a computer display.

For a SENTENCE consisting of a SUBJECT and PREDICATE: first run the procedure for the SUBJECT, insert a space, and then run the procedure for the PREDICATE.

For a SUBJECT consisting of a PRONOUN: just run the procedure for the PRONOUN.

For the PRONOUN *we*: insert *nous* into the output being generated.

For a PREDICATE consisting of a VERB: just run the procedure for the VERB.

For the VERB *know*: insert *connaiss*.

With this particular set of semantic procedures, the sentence “we know” will be translated to *nous* followed by a blank space followed by *connaiss*. You can easily verify this by starting with the semantic procedure for SENTENCE and tracing through the recursive cascade of procedure executions.

Each syntactic rule in a grammar must have a semantic procedure, even though the procedure might be quite trivial such as above when a SUBJECT is just a PRONOUN or a PREDICATE is just a VERB. This is because we need to make a connection at each level from SENTENCE all the way down to individual words like *we* and *know*. These connections give us a framework to extend our translation capabilities just by adding more syntactic rules plus their semantic procedures such as

PRONOUN->*they*

For the PRONOUN *they*: insert *ils*.

You may have noticed, however, our example above is incorrect. More so than English, French verbs must agree in person and number with their subject, and so the translation of *know* with the SUBJECT *we* should be *connaissons* (first person plural) instead of *connaiss* (the infinitive). Yet we cannot simply change the VERB semantic procedure above to “insert *connaissons*” because this would be wrong if the PRONOUN for SUBJECT becomes *they* (third person plural).

We need more elaborate semantic procedures here to get correct agreement. This is where various other PyElly commands have to come in; and in particular, we shall use local variables to pass information about number and person between the semantic procedures for our syntactic structures to govern their translations (see Section 6). Nevertheless, our PyElly framework of semantic procedures attached to each syntactic rule and called recursively will remain the same.

Semantic procedures must always be coded carefully for proper interaction and handling of details in all contexts. We would have to anticipate all the ways that constituent structures can come together in a sentence and provide for all the necessary communication between them. We can make the problem easier here by taking care to have lower-level structures be a part of only a few higher-level structures, but this will still require some advance planning.

Writing syntactic rules and their semantic procedures is actually a special kind of programming and will require programming skills. It will be harder than you first might think when you try to deal with natural languages like English or French. PyElly, however, is designed to help you to do this programming in a highly structured way, and it should be easier than trying to write the same kind of translation code explicitly in a language like PASCAL or BASIC or even LISP.



## 4. Defining Tables of PyElly Rules

By now, you should understand the idea of grammar rules and semantic procedures. This section will go into the mechanics of how you actually define them for PyElly in text files to be read in by PyElly at startup. To implement different applications such as translating English to French or rewriting natural language questions as structured data base queries, you just need to provide the appropriate files for PyElly to load.

The principal PyElly rules fall into five main types: (1) grammar, (2) vocabulary, (3) macro substitutions, (4) patterns for determining syntactic types, and (5) morphology. The grammar of a language for an application tends to reflect the capabilities supported by a target system, while a vocabulary tends to be geared toward a particular context of use; macros support particular users of a system, and special patterns tend to be specific to given applications. Separate tables of rules make it easier to tailor PyElly processing for different environments while allowing parts of language definitions to be reused.

This section will focus on the grammar, special pattern, and macro rule tables, which will be required by most PyElly applications. The creation and use of tables for vocabulary and morphology will be described in Section 9, “Advanced Programming: Vocabulary.” Some of the more technical details of semantic procedures for vocabulary definitions will also be postponed to Section 8, “Advanced Programming: Grammar.”

To make PyElly to something, you have set up an application defined by a specific set of language definition rules organized into tables. The current PyElly package defines each type of rule table as a Python class with an initialization procedure that reads in its rules from an external source, typically a text file. The names of the text input files associated with a particular application *A* should be named as follows:

`A.g.elly` for grammar rules and their semantic procedures.

`A.m.elly` for macro substitutions.

`A.p.elly` for special patterns.

You may replace the *A* here with whatever name you choose for your application, subject to the file-naming rules of the file system for your computer platform. Only the `A.g.elly` file is mandatory for any PyElly application; the other two may be omitted if you have no use for either substitutions or patterns. Section 7 will explain how PyElly will look for various language definition files for an application and read them in.

The rest of this section will describe the required formats of the definitions in the input files `A.g.elly`, `A.m.elly`, and `A.p.elly`. Normally you would create these files with a text editor or a word processor. The NotePad accessory on a Windows PC or TextEdit on a Mac will be quite adequate, although you may have to rename your files afterward because of their insistence on writing out files only with extensions like `.txt`.

An important element of most language rules will be syntactic structure names, seen in Section 2. These might also be called “syntactic types” or “parts of speech,” but they will

be more general than what we saw in grade school. The current implementation of PyElly in Python can handle up to 64 different structure types in its input files. Five of these structure types, however will be predefined by PyElly with special meanings.

SENT	Short for SENTENCE. Every grammar must have at least one rule of the form <code>SENT-&gt;X</code> or <code>SENT-&gt; X Y</code> . PyElly translation will always start by executing the semantic procedure for a <code>SENT</code> structure.
END	For internal purposes only. Avoid using it.
UNKN	Short for UNKNown. This structure type is automatically assigned to strings not known to PyElly through its various lookup options. (See Subsection 9.1 for more on this.)
...	For an arbitrary sequence of words in a sentence. This is for applications where much of the text input to process is unimportant. (See Section 8 for more details.)
PUNC	For punctuation. See Section 11.

You will of course have to make up your own names for any other structure types needed for a PyElly application. Names may be arbitrarily long in their number of characters, but may include only letters, digits, and periods (.); upper and lower case will be the same. You do not have to use traditional grammatical names like `NOUN`, but why be overly obscure here?

You may want to keep structure names unique in their first four characters. This is because PyElly may truncate names to that many characters in its formatted diagnostic output like parse trees (see Section 12, “PyElly Parsing”). The resulting tree might then be confusing if you have syntactic types like `NOUN` and `NOUNPHRASE`.

Here are some trivial, but functional, examples of grammar, macro, and pattern definition files:

```
# PyElly Definition File
# example.g.elly
g:sent->ss

—
g:ss->unkn

—
g:ss->ss unkn

—
```

```
# PyElly Definition File
# example.m.elly
i'm->i am
```

```
# PyElly Definition File
# example.p.elly
0 &# number 0 -1
```

These rules will be explained in separate subsections below.

## 4.1 Grammar (A.g.elly)

An A.g.elly text file may have four different types of definitions: (1) syntactic rules with their associated semantic procedures, (2) individual words with their associated semantic procedures, (3) general semantic subprocedures callable from elsewhere, and (4) initializations of global variable at startup. These definitions will be respectively identified in the A.g.elly file by special markers at the start of a line: G:, D:, P:, and I:. The definitions may be appear in any order.

### 4.1.1 Syntactic Rules

These must be entered as text in a strict line format. Syntactic rule definitions will follow the general outline as shown in monospaced font:

```
G:X->Y          # a marker + a syntax form
-               # a single <UNDERSCORE>,
                # omitted if no semantic
                # procedure follows
.               #
.               # the body of a generative
.               # semantic procedure
                #
—              # a double <UNDERSCORE>,
                # mandatory definition terminator
```

- In a G: line, PyElly will allow spaces anywhere except before an initial '#', within a syntactic structure name, or between the '-' and '>' of a rule.
- A '#' at the beginning of a line or a ' #' elsewhere indicates that a comment follows on the right; PyElly will ignore all comments within definition text.
- The same formatting applies for a PyElly syntactic rule of the form X->Y Z.
- The single and double underscores ( \_ or \_\_ ) here serve to mark the semantic procedures for a rule. A generative semantic procedure for a syntactic rule will always appear between such underscores.
- A cognitive semantic procedure will always appear before the single underscore, but this will be described later in Section 10.
- The actual basic actions for generative semantics will be described in Section 5 ("Operations for PyElly Generative Semantics").
- If a semantic procedure is omitted, various defaults apply; see Section 6 ("PyElly Programming Examples").

### 4.1.2 Grammar-Defined Words

In general, the vocabulary for a PyElly application should be separate from a grammar. For scalability, PyElly will store its vocabulary in an external Berkeley Database file; and Section 9 will describe how this should be set up. Some word definitions, however, may also appear alongside the syntactic rules in a grammar definition file. These will be called internal dictionary rules.

In particular, some words like THE, AND, and NOTWITHSTANDING are associated with a language in general instead of any particular content. These are probably best defined in a grammar file anyway. In other cases, there may also be so few words in a defined vocabulary that we may as well include them all internally in a grammar rather than externally.

The form of an internal word definition is similar to that for a grammatical rule:

```
D:w<-X          # a marker + a structure type X
                  #   + a word "w"
—                # a single <UNDERSCORE>
      .          #
      .          # a generative semantic procedure
      .          #
—                # a double <UNDERSCORE>,
                  #   mandatory definition terminator
```

- The `D:` is mandatory in order to distinguish a word rule from a grammatical rule of the form `X->Y`.
- The underscore separators are the same as for syntactic rules. A word definition may also have both cognitive and generative semantics.
- To suggest the familiar form of printed dictionaries, the word `w` being defined appears first, followed by its structure type `X` (i.e., part of speech). Note that the direction of the arrow `<-` is reversed from that of syntax rules.
- The `w` must be a single word, possibly hyphenated. Multi-word terms in an application must be defined in PyElly's external vocabulary or stitched together by grammar rules or macro substitution rules.

### 4.1.3 Generative Semantic Subprocedures

Every PyElly generative semantic procedure will be written in a special PyElly programming language for text manipulation. This language allows for named subprocedures, which need not be attached to a specific syntax rule or internal dictionary rule. Such subprocedures may be called by any generative semantic procedure for a PyElly rule or by another subprocedure. their definitions may appear anywhere in a \*.g.elly grammar file.

A subprocedure will take no arguments and return no values. All communication between semantic procedures must be through global or local variables and from the text passed in PyElly output buffers (see Section 5 for details). Calls to subprocedures may even be recursive, but if you do this, be careful to avoid infinite regression.

A subprocedure definition will have the following form:

```
P:n          # a marker + procedure name "n"
-           # a single <UNDERSCORE>,
            # mandatory
            #
            #
            # generative semantic procedure body
            #
            #
            # double <UNDERSCORE> delimiter,
            # mandatory
```

- Note the absence of any arrow, either  $\rightarrow$  or  $\leftarrow$ , in the first definition line.
- A procedure name  $n$  should be a string of alphanumeric characters without any spaces. It can be of any non-zero length. The case of letters is unimportant.
- The underscore separators are the same as for syntactic rules and word definitions, but they both will be mandatory for a subprocedure definition.
- A subprocedure definition may have only generative semantics. Cognitive semantics will not apply to a subprocedure and will always be ignored if specified.

### 4.1.4 Global Variable Initializations

PyElly global variables in a generative semantic procedure can be set in various ways. When such variables store important parameters referenced in a particular application grammar, it is helpful to be able to define them within the definition file for that grammar. In that way, the definition will be more readable and more easily maintained. The startup initialization of global variable  $x$  to the string  $s$  is accomplished by a  $\mathbb{I}:$  line in a grammar definition file:

```
 $\mathbb{I}:$   $x = s$ 
```

One must have one `I :` line for each global variable being initialized. Note that an `I :` line always stands by itself; there is no associated generative semantic procedure as in the case of `G :`, `D :`, and `P :` lines. An `I :` line may appear anywhere in a grammar definition file, but for clarity, it should be before any reference to it in a semantic procedure. For readability, you may freely put spaces around the variable name `x` and after the `=` sign here. For example,

```
I:iterate      = abcdefghijklm
I:joiner       = svnm
```

In the first initialization above, the `iterate` global variable is set to `abcdefghijklm`. A string value may have embedded space characters, but all leading and trailing spaces will be ignored and multiple consecutive embedded spaces will be collapsed to one.

## 4.2 Special Patterns (`a.p.elly`)

Many elements of text are too numerous to list out in a dictionary, but are recognizable by their form; for example, Social Security numbers, web addresses, or Russian surnames. PyElly allows you to identify such elements in input text by specifying the patterns that they conform to. That is how PyElly now deals with ordinary decimal numbers in text to be translated.

PyElly special patterns serve to assign a syntactic structure type (part of speech) to an otherwise unknown single word or other token in its input text. Any definition in a grammar's internal dictionary (see Subsection 4.1.2) or in its external vocabulary table (see Subsection 9.4) will override such pattern syntactic type assignment. For example, if `'123'` is a `NOUN` by a `D :` rule in a grammar table, then PyElly will no longer try to infer that it is a structural type `NUM`.

In general, we may have to compare multiple patterns in various ways to identify a particular text element. PyElly coordinates this kind of analysis with a finite-state automaton (FSA), which should be familiar to every aspiring computational linguist. This is not a physical machine, but a software algorithm working from a predefined set of rules telling it how to proceed step by step in matching up patterns from left to right in input text.

The key concept in an FSA is that of a state, which sums up how far along a text string the FSA has so far managed to match and what patterns it should look for next. An FSA will typically have multiple states, but one will always be a special starting state when the FSA is looking at the front of an input text with nothing yet matched.

At any given state, a PyElly FSA will have a list of patterns with possible wildcards to check against the input text at its current position. A wildcard here is a pattern element able to match against more than one text character; for example, any digit 0-9. PyElly wildcards are similar to the wildcards used in regular expressions, but are defined for natural language processing in particular; they will be defined in a listing below. This differs from the operation of FSA's seen elsewhere; they typically allow no wildcards.

Each pattern at a state will have associated actions to take upon any match. Usually, an FSA will move forward in its input string and go on to a next state, according to its predefined rules. There may be more than one such next state because a string at a given FSA state could match more than one pattern with wildcards. This is a complication, but everything is still equivalent to a regular FSA. It just makes rule sets more compact.

Some matches will have no next state in a PyElly FSA table, but instead specify a syntactic structure type. This tells a PyElly FSA that it has successfully completed a full match of input and can assign the given structure type to the portion of the input string matched so far. At this point, a normal FSA would be done, but PyElly will also have to examine all the matching possibilities arising from multiple next states for an FSA.

PyElly continues until all reachable states and patterns have been checked or until the FSA runs out of input. At that point, PyElly will return a positive match length if any final match has been made; 0, otherwise.

PyElly identifies each FSA current and next state by a unique non-negative integer, where the initial state is always 0. The absence of a next state after a match is indicated by -1. At each state, what to look for next is defined as a pattern of literal characters and wildcards. A \*.p.elly definition file will consist of separate lines each specifying a possible pattern for a given state, an optional PyElly syntactic structure type associated with any match, and a next state upon a match. These specifications comprise the table of rules that a PyElly FSA will work with.

Here is a simple PyElly file of some FSA pattern definitions:

```
# simple FSA to recognize syntactic structure types
# example.p.elly
#
# each input record is a 4-tuple
# STATE PATTERN SYNTAX NEXT

0 #, - 1
0 ##, - 1
0 ###, - 1
1 ###, - 1
1 ###$ NUM -1
0 &# - 2
2 . - 3
2 $ NUM -1
3 &#$ NUM -1
3 $ NUM -1
```

This recognizes entities of type NUM as plain integers like 1024, simple decimal values like 3.1416, and longer digit strings with commas like 1,001,053. A pattern line in general will have four parts as follows:

## PyElly User's Manual

state      Pattern      Syntactic Type      next

- The first part is an integer  $\geq 0$  representing a current PyElly automaton state.
- The second part is a pattern specifying arbitrary sequences of letters, numbers, and certain punctuation: hyphen (-), comma (,), period (.), slash (/). If these are present, they must be matched exactly within a word being analyzed.
- A pattern may have explicit characters and also various wildcards, which can match various substrings. Wildcards will be as follows:

#	will match a single digit 0 - 9
@	will match a single letter a - z or A - Z, possibly with diacritics
?	will match a single digit or letter
*	will match a n arbitrary sequence of non-blank characters, including a null sequence
&?	will match one or more letters or digits in a sequence
&#	will match one or more digits in a sequence
&@	will match one or more letters in a sequence
^	will match a single vowel
%	will match a single consonant
\$	will match the end of a word, but not add this to the extent of any matching

- Brackets [ and ] in a pattern will enclose an optional subsequence to match; only one level of bracketing is allowed and no wildcards are allowed inside.
- A pattern with a wildcard other than \$ by itself must always match at least one character; for example, [a] \* will be rejected.
- All final state patterns not ending with the \* or \$ wildcards will have a wildcard \$ appended automatically.
- The third part of a pattern line is a syntactic structure type (part of speech) like NOUN. A '-' here means that no type is specified.
- The fourth part is the next state to go to upon matching a specified pattern. This will be an integer  $\geq -1$ . A -1 state means a final state.

For example, the pattern ###-##-####\$ matches Social Security numbers, while the pattern (###) ###-####\$ matches a telephone number with an area code, which is not separated by a space. See Subsection 9.2.1 for more on possible number patterns. A wildcard character like & may be matched explicitly in a pattern by escaping it with a backslash character: \&.



### 4.3 Macro Substitutions (**A.m.elly**)

Macro substitution is a way of automatically replacing specific substrings in an input stream by other substrings. This is a useful capability to have in a language translator, and so PyElly includes it as yet another tool, adapting code from Kernighan and Plauger, *Software Tools*, Addison-Wesley, 1976.

The main difference in PyElly macro substitution versus *Software Tools* is that substrings to be replaced can be described with wildcards along with explicit characters to match and that substrings to replace parts of the original can include the parts of the original matching wildcards. Matching and replacement all happens before the start of PyElly parsing.

Macro substitution provides a convenient way of handling idioms, synonyms, abbreviations, and alternative spellings and of carrying out simple syntactic transformations awkward within the framework of context-free grammar rules. The general way to define a PyElly macro substitution rule is as follows:

P Q R → A B C D

- a. Each macro definition is limited to a single line. Since macros will be in their own \*.m.elly file, we need no identifier at the beginning of each line.
- b. The left and right sides of a substitution may have an arbitrary number of components, each being separated from the others by a blank.
- c. Upper and lower case is significant only on the right side.
- d. Input words matching the pattern of the left side will be replaced by the right side.
- e. The left side of a substitution rule may have patterns with wildcards; these will be the same as recognized in the special patterns described by the preceding subsection (3.2).
- f. A pattern may also have ' ' as a wildcard, which will match a single space character. This is not in wildcard list for the special patterns above because macro patterns can match multiple words, while special patterns can match only single words.
- g. A \\1, \\2, \\3, and so forth, on the right stands respectively for the first, second, third, and so forth, sequences of text matched by wildcard patterns on the left. Note that each of these applies to a sequence of contiguous wildcards; for example, the pattern ###abc@@@ on a match will associate the first 3 digits of a match with \\1 and the last three letters with \\2.
- h. When any macro is matched, its substitution will be done. Then all macros will be checked again against the modified result for other possible substitutions. When a macro eliminates its match entirely, though, further substitutions will be ended at that position.
- i. The order of macro definitions is significant. Those defined first in a definition file will always be applied first, possibly affecting the applicability of those defined afterward. Macros starting with a wildcard will always be checked after all others, however.

Macro substitutions are trickier to work with than grammatical rules because it is possible to define them to work at cross-purposes. You can even get into an infinite loop of substitutions if you are careless. Nevertheless, macros can greatly simplify a language definition when you use them properly and keep their patterns fairly short. They are

always applied before any syntactic analysis of a sentence with grammatical rules and so can override any tokenization rules in effect. Macros have no associated semantic procedures.

Here some examples of actual PyElly macro substitutions from its `texting` application:

```
*'ll -> \\1
percent* -> %
will not -> willnot
greater than or equal to -> >=
carry -ing -> with
receiv[e]* -> rcv\\1
#* @[.]m$ -> \\1\\2m
```

The last rule above will replace “10 p.m” with “10pm” to save space.

You often can use macro substitutions to handle idioms or other irregular forms that are exceptions to general language rules. They will always be applied after any inflectional stemming, but before any morphological stemming (see Subsection 9.1) with an unknown text element. Use them with care; PyElly will warn you about something possibly dangerous, but will not stop you from doing something catastrophic.

## 5. Operations for PyElly Generative Semantics

In Section 3, we saw examples of generative semantic procedures expressed in English. PyElly requires, however, that they be written in a special structured programming language for editing text in a series of output buffers. This language has conditional and iterative control structures, but generally operates at the nitty-gritty level of manipulating characters a few at a time.

Basically, PyElly generative semantics manages buffers and moves around text in them. The semantic procedures for various parts of a PyElly sentence all have to put their contributions for a translation into the right place at the right time. Proper coordination is critical; you have to plan everything out and control all interactions of procedures.

Every generative semantic procedure will be a sequence of simple commands, each consisting of an operation name possibly followed by arguments separated by blanks. These various operations are described below in separate subsections. For clarity, the operation names are always shown in uppercase here, but this does not matter to PyElly. Comments below begin with ' # ' and are not part of a command.

### 5.1 Insertion of Strings

These operations put a literal string at the end of the current PyElly output buffer:

APPEND any string	# put "any string" into current buffer
BLANK	# put a space character into buffer
SPACE	# same as BLANK
LINEFEED	# start a new line in current buffer
OBTAIN	# copy in the text for the first token # at the sentence position of the # phrase constituent being processed

### 5.2 Subroutine Linkage

For calling procedures of subconstituents for a phrase and returning from such calls:

LEFT	# calls the semantic procedure # for subconstituent structure # Y when a rule is of the form # X->Y or X->Y Z.
RIGHT	# calls the semantic procedure # for subconstituent structure # Z when a rule is of the form # X->Y Z.

```

RETURN                                # returns to caller

FAIL                                  # rejects the current parsing
                                     # of an input statement and
                                     # returns to the first place
                                     # where there is a choice of
                                     # of different parsings for
                                     # a constituent structure

```

### 5.3 Buffer Management

Processing starts with a single output text buffer. Spawning other buffers will often be helpful to keep the output of different semantic procedures separate for additional processing before their contents are finally joined together. You can put aside the current buffer and starting processing in a new buffer and then move text back and forth between the two buffers.

```

SPLIT                                # creates a new buffer and
                                     # directs processing to it

BACK                                  # redirects processing to end
                                     # of previous buffer while
                                     # preserving the new buffer

MERGE                                 # appends content of a new
                                     # buffer to the previous one,
                                     # deallocates the new one

```

These in effect allow a semantic procedure to be executed for its side effects without yet putting anything into the current output buffer. The `MERGE` operation can also be combined with string substitution:

```

MERGE /string1/string2/              # as above, except that all
                                     # occurrences of "string1"
                                     # in the new buffer will
                                     # be changed to "string2"

```

### 5.4 Local Variable Operations

Local variables can store a Unicode string. They are declared within the scope of a semantic procedure and will automatically disappear upon a return from the procedure.

```

VARIABLE x=string                    # declares variable x with
                                     # initial string value; if
                                     # no value is specified,
                                     # initialization is to null string

```

## PyElly User's Manual

```
SET x=string          # assigns string to local
                      # variable x of the most
                      # recent declaration
```

A string may contain any printing characters, but trailing spaces will be dropped. To handle single space characters specified by their ASCII names, you may use the following special forms:

```
VARIABLE x SP         # define variable x as single
                      # space char

SET x SP              # set variable x as single
                      # space char
```

Note the absence of the equal sign (=) here. PyElly will recognize SP, HT, LF, and CR as space characters here. This form can also be used with the IF, ELIF, WHILE, and BREAKIF semantic operations described below. You may write VAR as shorthand for VARIABLE; they are equivalent.

Some operations have a local variable as their second argument. These support assignment, concatenation of strings, and queuing.

```
ASSIGN  x=z           # assigns the value of local
                      # variable x to the local
                      # variable z in their most
                      # recent declarations

QUEUE   q=x           # appends the entire string stored
                      # in local variable x to any string
                      # stored already in local variable q

UNQUEUE x=q n         # removes the first n chars of the
                      # string stored in local variable
                      # q and assigns them to local
                      # variable x; if n is unspecified,
                      # the character count defaults to 1;
                      # if q has fewer than n chars, then
                      # x is just set to the value of q
                      # and q is set to the null string
```

The equal sign (=) must appear with SET and VARIABLE even when a second argument is missing; this is also required for UNQUEUE and QUEUE. If a lefthand local variable is undefined in a SET or ASSIGN operation, it will become automatically defined in the scope of the current generative semantic procedure.

## 5.5 Local Variable Set Operations

PyElly allows for manipulation of sets of strings, represented as their concatenation into a single string with commas between individual strings. For example, the set {"1","237","ab","uooo"} would be represented as the single string "1,237,ab,uooo". When local variables have been set to such list values, you can apply PyElly set-theoretic operations to them.

```

UNITE x<<z           # takes the union of the list values
                     # of local variables x and z
                     # and saves the result in x

INTERSECT x<<z       # intersects the list values
                     # of local variables x and z
                     # and saves the result in x

COMPLEMENT x<<z      # restricts the list values of
                     # of local variable x to those
                     # not in the list value for
                     # local variable z and saves
                     # the result in x

```

## 5.6 Global Variable Operations

Global variables are permanently allocated and are accessible to all semantic procedures through two restricted operations:

```

PUT x y              # store the value of local
                     # variable x in global
                     # variable y

GET x y              # the inverse of PUT

```

There is no limit on the total number of global variables. The global variables gp0, gp1, ... can be defined and set from a command line (see Section 7); you can define others yourself in generative semantic procedures by doing a PUT or a GET with a new global variable name.

## 5.7 Control Structures

Only two structures are supported: the IF-ELIF-ELSE conditional and the WHILE loop; they are as follows:

```

IF x=string          # if local variable x has
                     # value string, execute the
                     # following block of code

```

## PyElly User's Manual

```
ELIF x=string          # follows an IF; the test is
                        # made if all preceding
                        # tests failed and will
                        # control execution of
                        # following block of code
                        # (more than one ELIF can
                        # follow an IF)

ELSE                   # the alternative to take
                        # unconditionally after all
                        # preceding tests have failed

WHILE x=string          # the following block of code
                        # is repeatedly executed
                        # while the local variable
                        # x is equal to string

END                    # delimits a block of code and
                        # terminates an IF-ELIF-ELSE
                        # sequence or a WHILE loop
```

An **END** must terminate every **IF-ELIF-ELSE** sequence and every **WHILE** loop. PyElly will report a table definition error if any **END** is missing.

As in Subsection 5.4, we can check for single space characters here. For example,

```
IF x SP                # check if local variable x is
                        # a space character

ELIF x SP              #

WHILE x SP              #
```

Instead of **SP**, you may also have **HT**, **LR**, or **CR**.

A tilde (~) preceding the variable name **x** reverses the logical sense of comparison in all the checks above.

```
IF ~x=string           # test if x ≠ string
```

The **IF** and **ELIF** commands also have a form that allow for the testing a variable against a list of strings. PyElly allows for

```
IF   x=s, t, u         # test if x == s or x == t or x == u
ELIF x=s, t, u         # test if x == s or x == t or x == u
```

The strings to be compared against here must be separated by a comma (,) followed by a space. The space is essential for PyElly to recognize the listing here. The tests here can

be negated with a tilde (~) also. The checking of space characters as described above is not supported here.

Within a `WHILE` loop, you may also have

```
BREAK                # unconditionally break out
                     # of current WHILE loop

BREAKIF x=string      # if local variable x has
                     # value string, break out of
                     # current WHILE loop
```

The condition for `BREAKIF` can be negated with a tilde (~) as above. You can check for a single space character also.

## 5.8 Character Manipulation

These work with the current and next output buffers as indicated by < or > in a command; `x` specifies a source or target local variable to work with.

```
EXTRACT > x n        # drops the last n chars of
                     # the current output buffer and
                     # sets local variable x to the
                     # string of dropped characters

EXTRACT x < n        # drops the first n chars of
                     # the next output buffer and
                     # sets local variable x to the
                     # string of dropped characters

INSERT < x            # insert the chars of local
                     # variable x to the end of the
                     # current output buffer

INSERT x >           # insert the chars of local
                     # variable x to the start of the
                     # next output buffer

PEEK x <             # get a single char from
                     # start of next output buffer
                     # without removing it

PEEK > x             # get a single char from
                     # end of current output buffer
                     # without removing it

DELETE n <           # deletes n chars from the
                     # start of the next output buffer
```



## PyElly User's Manual

```
DELETE n >          # deletes n chars from the
                    # end of the current output
                    # buffer

STORE x k            # save last deletion in a current
                    # procedure in local variable
                    # except for last k chars when
                    # k > 0 or the first k chars
                    # when k < 0; if unspecified,
                    # k defaults to 0

SHIFT n <           # shifts n chars from
                    # the start of the next output
                    # buffer to the end of the
                    # current output buffer

SHIFT n >           # shifts n chars from
                    # the end of the current output
                    # buffer to the start of the
                    # next output buffer
```

If *n* is omitted for the `EXTRACT` operation above, it is assumed to be 1. If the `<` or `>` are omitted from a `DELETE` or a `SHIFT`, then `<` is assumed. All the characters removed by any form of `DELETE` can be accessed by `STORE`.

The `DELETE` operation also has three variants

```
DELETE >            # this deletes every char
                    # in the current buffer

DELETE <            # this deletes every char
                    # in the next buffer

DELETE FROM s        # this deletes an indefinite
                    # number of chars starting from
                    # the string s in the current
                    # buffer up to the end

DELETE TO s          # this deletes an indefinite
                    # number of chars up to and
                    # including the string s
                    # in the next buffer
```

If the argument *s* is omitted for `DELETE FROM` or `DELETE TO`, it is taken to be the string consisting of a single space character. If *s* is not found in the current or the next buffer for `DELETE FROM` or `DELETE TO`, all of that buffer will be deleted. As with the regular `DELETE` operation, any characters removed by this command can be accessed by the `STORE` command.

## 5.9 Selection From a Table

This operation that uses the value of a local variable to select a string for appending at the end of the current output buffer. It has the form

```
PICK x table           # select from table according
                       # to the value of x
```

The table argument is a literal string of the form

```
(v1=s1#v2=s2#v3=s3#...vn=sn#)
```

If the value of local variable `x` here is equal to substring `vi`, then substring `si` will be inserted. If there is no match, nothing will be inserted, but when `vn` is null, then `sn` will be inserted if the variable `x` matches no other `vi`.

For example, the particular `PICK` operation

```
PICK x (uu=aaaa#vv=bbbb#ww=cccc#=dddd#)
```

in a generative semantic procedure is equivalent to the code

```
IF    x=uu
  APPEND aaaa
ELIF  x=vv
  APPEND bbbb
ELIF  x=ww
  APPEND cccc
ELSE
  APPEND dddd
END
```

but the `IF-ELSE` form is much less compact. A `PICK` table will be saved in a generative semantic procedure as a Python hash object for faster lookup.

The operation

```
PICK x (=dddd#)
```

will append `dddd` for any `x`, including `x` being set to the null string.

## 5.10 Buffer Searching

There is one search operation in forward and reverse forms. These assume existence of a current and a new buffer as the result of executing `SPLIT` and `BACK`.

```
FIND s >                                # the contents of the new
                                         # buffer will be shifted to the
                                         # current buffer up to the first
                                         # occurrence of string s

FIND s <                                # as above, but transferring
                                         # is in the other direction
```

If no substring `s` is given, an entire buffer will be moved. If the `<` or `>` is omitted, then `>` is assumed.

## 5.11 Execution Monitoring

To track the execution of semantic procedures when debugging them, you can use the command:

```
TRACE                                # show processing status in tree
```

When executed in the semantic procedure for a phrase, this will print to the standard error stream the starting token position of the phrase in a sentence, its syntax type, the index number of the syntactic rule, and the degree of branching of the rule. Within a named subprocedure, PyElly will follow the chain of calls back to the first semantic procedure for an actual phrase. Here is some `TRACE` output:

```
TRACE @0 type=field rule=127 (1-br) stk=9 buf=1 (2 chars)
```

This specifies the current token position (=0 here) of a phrase, its syntactic type and PyElly rule index, the type of rule, the depth of the calling stack for generative semantic procedures, the total buffer count, and the number of characters in the current output buffer. If a subprocedure named `pn` (see Subsection 5.13) has called the current generative semantic procedure either directly or indirectly, then this will be reported also. The output above then becomes

```
TRACE @0 type=field rule=127 (1-br) stk=9 in (pn) buf=1 (2 chars)
```

If there are multiple named subprocedures in the chain of calls for the current generative semantic procedure, then only the most recent is reported here.

To show the current string value of a local variable `x`, you can insert this command into a semantic procedure:

```
SHOW x message. . . . # show value of local variable
```

This writes the ID number of the phrase being interpreted, the name of the variable in its generative semantic procedure, the current string value of the variable, and an optional message to the standard error stream to help in debugging. For example,

```
SHOW @phr n : [message. . . ] VAR x= [012345]
```

You can omit the message string here.

## 5.12 Capitalization

PyElly has only two commands to handle upper and lower case in output.

```
CAPITALIZE          # capitalize the first char
                    # in the next buffer after a
                    # split and back operation
```

This operates only on the next output buffer. If you fail to do a `SPLIT` and `BACK` operation to create a next output buffer before running this command, you will get a null pointer exception, which will halt PyElly.

```
UNCAPITALIZE        # uncapitalize the first char
                    # in the next buffer after a
                    # split and back operation
```

The restrictions for `CAPITALIZE` apply here also.

## 5.13 Semantic Subprocedure Invocation

If `DO` is the name of a semantic subprocedure defined with `P :` in a PyElly grammar table, then it can be called in a generative semantic procedure by giving the name in parentheses:

```
(DO)                # call the procedure called DO
```

The subprocedure name must be defined somewhere in a PyElly `A.g.elly` file. This definition does not have to come before the call. When a subprocedure finishes running, execution will return to the point just after where it was called.

A subprocedure call will always take no arguments. If you want to pass parameters, you must do so through local or global variables or in a buffer. Results from a subprocedure will have to be returned in the same way.

The null subprocedure call `()` with no name is always defined; it is equivalent to a generative semantic procedure consisting of just a `RETURN`. This is normally used only for PyElly vocabulary definitions with no associated semantics (see Subsection 9.4).

## 6. Simple PyElly Programming Examples

We are now ready to look at some simple examples of semantic procedures for PyElly syntax rules, employing the mechanisms and operations defined in the preceding sections. Sections 7 and 8 will discuss more advanced capabilities.

### 6.1 Default Semantic Procedures

The notes in the Section 4.1.1 of this manual mentioned that omitting the semantic procedure for a syntax rule would result in a default procedure being assigned to it. Now we can finally define those default procedures. A rule of the form  $G : X \rightarrow Y \ Z$  will have the default

```
— LEFT
  RIGHT
```

Note that a `RETURN` command is unnecessary here as it is implicit upon reaching the end of the procedure. You can always put one in yourself, however.

A rule of the form  $G : X \rightarrow Y$  has the default semantic procedure

```
— LEFT
```

A rule of the form  $D : w \leftarrow X$  has the default

```
— OBTAIN
```

These are automatically defined by PyElly as subprocedures without names. They do nothing except to implement the calls and returns needed minimally to maintain communication between the semantic procedures for the syntactic rules associated with the structure of a sentence derived by a PyElly analysis.

In the first example of a default semantic procedure above, a call to the procedure for the left constituent structure  $X$  comes first, followed immediately by a call to the procedure for the right constituent  $Y$ . If you wanted instead to call the right constituent first, then you would have to supply your own explicit semantic procedure, writing

```
— RIGHT
  LEFT
```

In the second example above of a default semantic procedure, there is only one constituent in the syntactic rule, and this will always be a left constituent by convention. We have no right constituent, and trying to call its semantic procedure with a `RIGHT` command would result in an error and halt PyElly.

In the third example of a default semantic procedure above, which defines a grammatical word, there is neither a left nor a right constituent; and so we will execute an `OBTAIN`. Either a `LEFT` or a `RIGHT` command here would result in an error.

## 6.2 A Simple Grammar Example

We now give an example of a nontrivial PyElly grammar. The problem of making subjects and predicates agree in French was discussed previously in Section 3. Here we make a start at a solution by handling *elison* and the present tense of first conjugation verbs in French plus the irregular verb `AVOIR` “to have.” For the relationship between a subject and a predicate in the simplest possible sentence, we have the following syntactic rule plus semantic procedure.

```
G:SENT->SUBJ PRED
—
VAR PERSON=3      # can be 1, 2, or 3
VAR NUMBER=s      # singular or plural
LEFT              # for subject
SPLIT
RIGHT             # for predicate
BACK
IF PERSON=1
  IF NUMBER=s
    EXTRACT X <   # letter at start of predicate
    IF X=a, e, è, é, i, o, u
      DELETE 1    # elison j'
      APPEND `    #
    ELSE
      BLANK       # otherwise, predicate is separate
    END
    INSERT < X    # put predicate letter back
  END
ELSE
  BLANK          # predicate is separate
END
ELSE
  BLANK          # predicate is separate
END
MERGE            # combine subject and predicate
APPEND !
—
```

The two local variables `NUMBER` and `PERSON` are for communication between the semantic procedures for `SUBJ` and `PRED`; they are set by default to “singular” and “third person”. The semantic procedure for `SUBJ` is called first with `LEFT`; then the semantic procedure for `PRED` is called with `RIGHT`, but with its output in a separate buffer. This lets us adjust the results of the two procedures before we actually merge them; here the commands in the conditional `IF-ELSE` clauses are to handle a special case of *elison* in French when the subject is first person singular and the verb begins with a vowel.

```
G:SUBJ->PRON
```

---

The above rule allows a subject to be a pronoun. The default semantic procedure for a syntactic rule of the form `X->Y` as described above applies here, since none is supplied explicitly.

```
D:i<-PRON
```

---

```
  APPEND je
  SET PERSON=1
```

---

```
D:you<-PRON
```

---

```
  APPEND vous
  SET PERSON=2
  SET NUMBER=p
```

---

```
D:it<-PRON
```

---

```
  APPEND il
```

---

```
D:we<-PRON
```

---

```
  APPEND nous
  SET PERSON=1
  SET NUMBER=p
```

---

```
D:they<-PRON
```

---

```
  APPEND ils
  SET NUMBER=p
```

---

These internal dictionary syntax rules define a few of the personal pronouns in English for translation. The semantic procedure for each rule appends the French equivalent of a pronoun and sets the `PERSON` and `NUMBER` local variables appropriately. Note that, if the defaults values for these variables apply, we can omit an explicit `SET`.

Continuing, we fill out the syntactic rules for our grammar.

```
G:PRED->VERB
```

---

This defines a single `VERB` as a possible `PRED`; the default semantic procedure applies again, since no procedure is supplied explicitly here.

Now we are going to define two subprocedures needed for the semantic procedures of our selection of French verbs.

```
P:default
```

```
— PICK PERSON (1=ons#2=ez#3=ent#)
```

---

```
P:1cnjg
```

```
— IF NUMBER=s
    PICK PERSON (1=e#2=es#3=e#)
ELSE
    (default)
END
```

---

Semantic subprocedures `default` and `1cnjg` choose an inflectional ending for the present tense of French verbs. The first applies to most verbs; the second, to first conjugation verbs only. We need to call them in several places below and so define the subprocedures just once for economy and clarity.

```
D:sing<-VERB
```

```
— APPEND chant      # root of verb
  (1cnjg)           # for first conjugation inflection
```

---

```
D:have<-VERB
```

```
— IF NUMBER=s
    PICK PERSON (1=ai#2=ais#3=a#)
ELSE
    IF PERSON=3
        APPEND ont  # 3rd person plural is irregular
    ELSE
        APPEND av   # 1st and 2nd person are regular
        (default)
    END
END
```

---

We are defining only two verbs to translate here. Other regular French verbs of the first conjugation can be added by following the example above for “sing”. Their semantic



procedures will all append their respective French roots to the current output buffer and call the subprocedure `lcnjg`.

The verb AVOIR is more difficult to handle because it is irregular in most of its present tense forms, and so its semantic procedure must check for many special cases. Every irregular verb must have its own special semantic procedure, but there are usually only a few dozen such verbs in any natural language.

Here is how PyElly will actually process input text with this simple grammar. The English text typed in for translation is shown in uppercase on one line, and the PyElly translation in French is shown in lowercase on the next line.

**YOU SING**  
**vous chantez!**

**THEY SING**  
**ils chantent!**

**I HAVE**  
**j'ai!**

**WE HAVE**  
**nous avons!**

**THEY HAVE**  
**ils ont!**

The example of course is extremely limited as translations go. For full-scale translation, we would also take English inflectional stemming into account, use macro substitutions to take care of irregularities on the English side like `has`, and handle other subtleties. We also have various tenses other than present. You should, however, be able to envision now what a full PyElly grammar should look like; it will take much more work to complete, but would be a straight extension of what we have seen above.

## 7. Running PyElly From a Command Line

We have so far described how to set up definition text files to create the various tables to guide PyElly operation. This section will show you how to run PyElly for actual language analysis, but first we will have to take care of some preliminary setup. That should be fairly straightforward, but computer novices may want to get some technical help here.

To begin with PyElly was written entirely in version 2.7 Python, which seems to be the most widely preinstalled by computer operating systems. The latest version of Python is 3.\*, but unfortunately, this is incompatible with 2.7. So make sure you have the right version here. Python is free software, and you can download a 2.7 release from the Web, if necessary. The details for doing so will depend on your computing platform.

After getting Python, you should next get Berkeley Database, a free database package used by PyElly to handle its vocabulary rules. Although you can get by without Berkeley Database, this will severely limit what you can do in PyElly. Get the latest release, currently 6.1.19, available from the Oracle website for most operating systems.

There is an additional complication with Berkeley Database. Python originally built in support for this, but that was dropped in version 2.7. To use Berkeley Database in Python, you must also download and install the free third-party `bsddb3` module. Again, details here will depend on your computing platform (see Appendix C). To save grief, go through a software management tool here instead of trying to do all the installations by individual commands from your keyboard.

Once you have Python with `bsddb3` and Berkeley Database installed, you are ready to download the full PyElly package from GitHub. This is open-source software under a BSD license, which means that you can do anything you want with PyElly as long as you identify in your own documentation where you got it. All PyElly Python source code is free, but under copyright.

The Python code making up PyElly currently consists of 58 modules comprising over 16,000 source lines altogether. A beginning PyElly user really needs to be familiar with only three of the modules.

**`ellyConfiguration.py`** - defines the default environment for PyElly processing. Edit this file to customize PyElly to your own needs. Most of the time, you can leave this module alone.

**`ellyBase.py`** - sets up and manages the processing of individual sentences from standard input. You can run this for testing or make it your programming interface if you want to embed PyElly in a larger application.

**`ellyMain.py`** - runs PyElly from a command line. This is built on top of `EllyBase` and is set up to extract individual sentences from continuous text in standard input.

The `ellyBase` module reads in `*.*.elly` language definition files to generate the various tables to guide PyElly analysis of input data. Section 4 introduced three of them. For a

given application `A`, these were `A.g.elly`, `A.p.elly`, and `A.m.elly`, with only the `A.g.elly` file mandatory. Subsequent sections of this user manual will describe the other `*.*.elly` definition files.

The PyElly tables created for an application `A` will be automatically saved in two files: `A.rules.elly.bin` and `A.vocabulary.elly.bin`. The first is a Python pickled file, which is not really binary since you can look at it with a text editor, but this will be hard for people to read. The second is a binary database file produced by Berkeley Database from definitions in a given `A.v.elly` (see Subsection 9.4 for an explanation).

If the `*.elly.bin` files exist, `ellyBase` will compare their creation dates with the modification dates of corresponding `*.*.elly` definition files and create new tables only if one or more definition files have changed. Otherwise, the existing PyElly language rule tables will be reloaded from the `*.elly.bin` files.

The files `*.rules.elly.bin` keep track of which version of PyElly they were created under. If this does not agree with the current version of PyElly, then PyElly will immediately exit with an error message that the rule file is inconsistent. To proceed, you must then delete all of your `*.elly.bin` files so that they can be regenerated from your latest language definition files.

In most cases, `ellyBase` will try to substitute a file `default.x.elly` if an `A.x.elly` file is missing. This may not be what you want. You can override this behavior just by creating an empty `A.x.elly` file. The standard PyElly download package includes six examples of definition files for simple applications to show you how to set everything up (see Section 14).

You can see what `ellyBase` does by running it directly with the command line:

```
python ellyBase.py [name]
```

This will then generate the PyElly tables for the `name` application and prompt for individual sentences to analyze with them. If no application is specified in the command line, the default is `test`. Running `ellyBase` directly can be helpful in tracking down problems in a language definition because it will provide a detailed dump of grammar rules and closer monitoring of PyElly analysis. It also lets you prebuild `*.elly.bin` files so that PyElly can subsequently start up faster.

For an application with batch processing of input sentences not necessarily on separate lines, you normally will invoke `ellyMain` from a command line. The `ellyMain.py` file is a straight Python script that reads in general text and allows you to specify various options for PyElly language processing. Its full command line is as follows in usual Unix or Linux documentation format:

```
python ellyMain.py [ -b ][ -d n ][ -g v0,v1,v2,... ][ -p ][ -noLang ] [name] < text
```

## PyElly User's Manual

where `name` is an application identifier like `A` above and `text` is an input source for PyElly to translate. If the identifier is omitted, the application defaults to `test`.

The commandline flags here are all optional. They will have the following interpretations in PyElly `ellyMain`:

<code>-b</code>	operate in batch mode with no prompting; PyElly will otherwise run in interactive mode with prompting when its text input comes from a user terminal.
<code>-d n</code>	set the maximum depth for showing a PyElly parse tree to an integer <code>n</code> . This can be helpful when input sentences are quite long, and you do not want to see a full PyElly parse tree. Set <code>n = 0</code> to disable parse trees completely. See Section 12 for more details.
<code>-g v0, v1, v2, ...</code>	define the PyElly global variables <code>gp0</code> , <code>gp1</code> , <code>gp2</code> , ... for PyElly semantic procedures with the respective specified string values <code>v0</code> , <code>v1</code> , <code>v2</code> , ...
<code>-p</code>	show cognitive semantic plausibility scores along with translated output. If semantic concepts are defined, PyElly will also give the contextual concept of the last disambiguation according to the order of interpretation by generative semantics. This is intended mainly for debugging, but may be of use in some applications (see <code>disambig</code> , described in Section 14).
<code>-noLang</code>	do not assume that input text will be in English; the main effect is to turn off English inflectional stemming (See Section 11).

When `ellyMain` starts up in interactive mode, you will see the following message:

```
PyElly v1.0, Natural Language Filtering  
Copyright 2014 under BSD open-source license by C.P. Mah  
All rights reserved
```

```
reading <a> definitions  
recompiling language rules
```

```
Enter text with one or more sentences per line.  
End input with E-O-F character on its own line.
```

```
>>
```

You may now enter multiline text at the `>>` prompt. PyElly will process this exactly as it would handle text from a file or a pipe. Sentences can extend over several lines, or a single line can contain several sentences. PyElly will automatically find the sentence boundaries according to its current rules and divide up the text for analysis.

As soon as PyElly reads in a full sentence, it will try to write a translation to its output. In interactive mode, this will be after the first linefeed after the sentence because PyElly has to read a full line before it can proceed. Linefeeds will NOT break out of the `ellyMain` input processing loop, although two consecutive linefeeds will terminate a sentence even when punctuation is absent. End your input with an EOF (control-D on Unix and Linux,

control-Z in Windows). A keyboard interrupt (control-C) will break out of `ellyMain` with no further processing.

PyElly `*.*.elly` language definition files should be in UTF-8 encoding and may contain arbitrary Unicode except for grammar symbol names. As text input to translate, however, PyElly currently accepts only ASCII and Latin-1 letters; all other input characters will be treated as spaces. The `chinese` application described in Section 14 uses definition files with both traditional and simplified Chinese characters in UTF-8.

All PyElly translation output will be UTF-8 characters written to the standard output stream, which you may redirect to save to a file or pipe to other modules outside of PyElly. PyElly parse trees and error messages will also be in UTF-8 and will go to the standard error stream, which you can also redirect. Historically, the predecessors of PyElly have always been filters, which in Unix terminology means a program that reads from standard input and writes a translation to standard output.

Here is an example of interactive PyElly translation with a minimal set of language rules (`echo.*.elly`) defining a simple echoing application:

```
>> Who gets the gnocchi?

=[who get -s the gnocchi?]
```

where the second line is actual output from `ellyMain`. Elly by default will convert upper case to lower, and will strip off English inflectional endings as well as `-ER` and `-EST`. You can get stricter echoing by turning off inflectional stemming and morphological analysis.

By default, PyElly will look for the definition files for an application in your current working directory. You change this by editing the value for the symbol `baseSource` in `ellyConfiguration.py`. The various PyElly applications described in Section 14 are distributed in the `applcn` subdirectory under the main directory of Python source files, resources, and documentation.

PyElly `*.py` modules by default should be in your working directory, too. You can change where to look for them, but that involves resetting environment variables for Python itself. PyElly is written as separate Python modules to be found and linked up whenever you start up PyElly. This is in contrast to other programming languages where modules can be prelinked in a few executable files or packaged libraries.

On the whole, you must be comfortable with computing at the at the level of command lines in order to run PyElly either by `ellyMain.py` or by `ellyBase.py`. There is as yet no graphical user interface for PyElly. The current PyElly implementation may be a challenge to computer novices unfamiliar with Python or with commandline invocation.

## 8. Advanced Capabilities: Grammar

As noted above, PyElly language analysis is built around a parser for context-free languages to take advantage of extensive technology developed for parsing computer programming languages. So far, we have stayed largely in the context-free domain except for macro substitution prior to parsing and use of local variables shared by generative semantic procedures to control translation.

You can actually accomplish a great deal with such basics alone, but for more challenging language analysis, PyElly supports other capabilities beyond the confines of pure context-free languages. These include extensions to grammar rules like syntactic and semantic features and the special . . . syntactic type mentioned earlier. Other extensions related to vocabularies are covered in the next section.

The handling of sentences and punctuation in continuous text is also normally a topic of grammar, but PyElly breaks this out as a separate level of processing for modularity. The details on this will be discussed in Section 11.

### 8.1 Syntactic Features

PyElly currently allows for only 64 distinctive syntactic types, including predefined types like `SENT` and `UNKN`. If needed, you get more types by redefining the variable `NMAX` in the PyElly file `grammarTable.py`, but there is a more convenient option here. PyElly also lets you qualify syntactic types by syntactic features, which in effect greatly multiplies the total number of syntactic types available.

Syntactic features are binary tags serving to subcategorize syntactic types, appearing in Noam Chomsky's seminal work *Syntactic Structures* (1957). Currently, PyElly allows up to sixteen syntactic features for each class of syntactic types. You can define the classes and name the features however you want. If really needed, you can get more than sixteen syntactic features by redefining the variable `FMAX` in `grammarTable.py`.

The advantage of syntactic features is that grammar rules can disregard them. For example, a `DEFINITE` syntactic feature would allow definite noun phrases to be identified in a grammar rule without having to introduce a new structural type like `DNP`. Instead, we would have something like `NP[:DEFINITE]`. A grammar syntax rule like `PRED->VP NP` would apply to `NP[:DEFINITE]` as well as to plain `NP`. With a new syntax type like `DNP`, we would also have to add the rule `PRED->VP DNP`.

PyElly syntactic features are expressed by an optional bracketed qualifier appended to a syntactic structural type specified in a rule. The qualifier takes the form

`[oF1, F2, F3, ..., Fn]`

where “o” is a single-character identifier for a set of features for a specific class of syntactic types and `F1, ..., Fn` are feature names composed of alphanumeric characters, possibly preceded by a prefix ‘-’ or ‘+’.

Allowing multiple sets of feature names is for convenience only. Each set will have to refer to the same `FMAX` feature bits defined for a phrase node in a PyElly parse tree. So if you want to define multiple name sets, then make sure that their usage is consistent. Especially avoid a syntactic type occurring with syntactic feature names from more than one set because features with different names then may not be unique.

Syntactic features in language rules must immediately follow a syntactic type name with no spaces in between. PyElly allows spaces in the listing of individual syntactic feature names for easier reading, but any space before the starting left bracket ( `[` ) will always be rejected as an error.

A syntax rule with feature names might appear as follows:

```
G:NP[:DEFINITE,*RIGHT]->THE NP[:~DEFINITE,*RIGHT]
```

This specifies a rule of the form `NP->THE NP`, but with additional restrictions on applicability. The `NP` as specified on the right side of the rule must have the feature `*RIGHT`, but not `DEFINITE`. If the condition is met, then the resulting `NP` structure as specified on the left of the rule is defined with the features `DEFINITE` and `*RIGHT`. The `‘:’` is the feature class identifier here for the `DEFINITE` and `*RIGHT` features.

For clarity, identifiers should be a punctuation character other than brackets or `‘+’` or `‘-’`. You should have at most five or six sets of syntactic features, although PyElly sets no upper limit here. Just remember that syntactic features are supposed to simplify grammars, not make them more complicated.

The special feature name `*RIGHT` (or equivalently `*R`) will be defined automatically for all syntactic feature sets. Setting this feature on the left side of a syntactic rule will have the side effect of making that constituent structure inherit any and all syntactic features of its rightmost immediate subconstituent as specified in the rule. This provides a convenient mechanism for passing syntactic features up a parse tree without having to say what exactly they are.

The special feature name `*LEFT` (or equivalently `*L`) will also be automatically defined. This will work like `*RIGHT`, except that inheritance will be from the leftmost immediate subconstituent. You can specify both `*LEFT` and `*RIGHT` in the features for a syntax type, but usually just one will suffice.

A third special feature name `*UNIQUE` will be predefined for all PyElly syntactic feature sets. It has the sole purpose of preventing a phrase from matching any other phrase in PyElly ambiguity checking while parsing. This and the other special syntactic features may not be redefined, but they will be counted in the number of syntactic features available for any given grammar.

## 8.2 The . . . Syntactic Type

When the . . . type shows up in a grammar, PyElly automatically defines a syntax rule that allows phrases to be empty. If you could write it out, the rule would take the form

```
. . .->
```

This is sometimes called a zero rule, which PyElly will not allow you to specify explicitly in a \*.g.elly file for any syntactic type on the left. In strict context-free grammars, any rule having a syntactic structural type going to an empty phrase is forbidden. Such rules are allowed only in so-called type 0 grammars, the most unrestricted of all; but the languages described by such grammars tend to be avoided because of the difficulty in parsing them.

With . . . as a special syntactic type, however, PyElly achieves much of the power of type 0 grammars without giving up the parsing advantages of context-free grammars. The advantage with . . . in PyElly is that it allows a grammar to be more compact when this syntactic type is applicable. For example, suppose that we have the rules

```
z->x a
z->x b
z->x c
z->x d
z->a
z->b
z->c
z->c
x->unkn
x->x unkn
```

where `unkn` is the predefined PyElly from Section 4 (this will be explained more fully in Section 9.1). Now if `x` is not of interest to us in an eventual translation, then we can replace all the above with just the rules

```
z->. . . a
z->. . . b
z->. . . c
z->. . . d
. . .->unkn
. . .->. . . unkn
```

The . . . type was intended specifically to support keyword parsing, which recognizes a limited number of words in input text and more or less ignores anything else. A PyElly grammar to support such parsing can always be written without . . ., but would be unwieldy. The `doctor` application for PyElly illustrates how this kind of keyword grammar would be set up; it includes syntax rules like the following:



```

g:ss->x ...
—
g:x[@*right]-> ... key
—
g:...->unkn ...

```

The syntactic type `key` here represents all the various kinds of key phrases to recognize in a psychiatric dialog; for example, “mother” and “dream”. We can get away with only one syntactic type here because, with about a dozen syntactic features available for it, we can distinguish between 215 different kinds of key phrases.

The actual responses of our script will be produced by semantic procedures for the rules defining `x[@*right]` phrases. Note that different responses to the same keyword must be listed as separate rules with the same syntactic category and features. A simplified listing of grammar rules here might be

```

g:sent[@*right]->ss
—
g:x->... key
—
g:key[@ 0,1]->fmly
—
g:ss[.*right]->x[@ 0, 1,-2,-3,-4,-5,-6] ...
—
  append TELL ME MORE ABOUT YOUR FAMILY
—
g:ss[.*right]->x[@ 0, 1,-2,-3,-4,-5,-6] ...
—
  append WHO ELSE IN YOUR FAMILY
—
d:mother <- fmly
—
g:...->unkn
—
g:...->unkn ...
—

```

This defines two different possible responses for `key[@ 0,1]` in our input. PyElly ambiguity handling will then automatically alternate between them (see Section 10).

The grammar here is incomplete, recognizing only sentences with a single keyword and nothing else. To allow for sentences without a keyword, we also need a rule like

```

g:ss->...
—

```

The `...` syntactic type also has the restriction that you cannot specify syntactic features for it. If you put something like `... [ .F1, F2, F3 ]` in a PyElly rule, it be treated as just `...`. This is to help out the PyElly parser, which is already working hard enough.

PyElly will also block you from defining a rule like

```
g:...->...
```

—

or like

```
g:X->... ..
```

—

where  $x$  is any PyElly syntactic type, including . . . .

The . . . syntactic type can be quite tricky to use effectively in a language description, but it is even trickier for PyElly to handle as an extension to its basic context-free parsing. The various restrictions here are a reasonable compromise to let us do what we really need to do. See Subsection 12.3.3 for details on how PyElly parsing actually handles . . . .

## 9. Advanced Capabilities: Vocabulary

PyElly operates by reading in, analyzing, and rewriting out sentences. To succeed here, it requires syntactic and semantic information for every text element that it encounters: words, names, numbers, identifiers, punctuation, and so forth. Certain text elements like punctuation will be fairly limited, but defining all the rest can be a big undertaking even for some fairly simple applications.

In all our PyElly examples so far here, we have already seen several ways of defining text elements.

- An explicit `D:` rule in a grammar.
- Assignment of syntactic information through matching of specified patterns.
- Making use of the predefined `UNKN` syntactic type.

These are fine with small vocabularies, but useful natural language applications must deal with hundreds or even tens of thousands of distinct terms. These may not fall into obvious patterns; and stuffing them all into a `*.g.elly` grammar file will demand more keyboard entry than most people care to do. Treating most text elements as `UNKN` is always a fallback option, but this essentially is giving up.

There is no perfect solution here. PyElly can only try to provide a user enough vocabulary definition options to make the overall task a little less painful. So, in addition to the methods above, PyElly also incorporates builtin analysis of unknown words to infer a syntactic type, plug-in code for recognizing complex entities like numbers, time, and dates, and vocabulary tables loaded from external databases. These will be described in separate subsections below, but as background, we first need to explain better how the `UNKN` syntactic type works.

### 9.1 Working with the UNKN Syntactic Type

We have run across the `UNKN` syntactic type several times already in this manual. Whenever text element `xxxx` in its input cannot be otherwise identified by PyElly, it will be assigned the type `UNKN`. In effect, PyElly generates a temporary rule of the form:

```
D:xxxx <- UNKN
—
OBTAIN
—
```

Such a rule is in effect only while PyElly is processing the current input sentence.

By itself, `UNKN` solves nothing. It just gives PyElly a way of working with unknown elements, and you still are responsible for supplying the grammar rules and associated semantics to tell PyElly how to interpret a sentence having `UNKN` as one of its subconstituents. The simplest possibility here is to make some guesses; for example,

G:NOUN->UNKN

---

G:VERB->UNKN

---

These two rules allow an unknown word to be treated as either a noun or a verb. So, when given a sentence containing unknown `xxxx`, PyElly can try to analyze it with both of its possible syntactic types. If only one results in a successful parse, then we have managed to get past the problem of having no definition for `xxxx`. If neither works out, we have lost nothing; if both work out, then PyElly can try to figure out which is the more plausible using the cognitive semantic facilities described in Section 10.

An unknown word can also be resolved by looking at how it is put together. For example, the word UNREALIZABLE may be missing from a vocabulary, but it could be broken down into UN+ +REAL -IZE -ABLE, allowing us to identify it as an adjective based on the root word REAL, which is much more likely to be defined already in a vocabulary. PyElly develops this idea will be further, and this will be described in the immediately following subsections.

## 9.2 Breaking Down Unknown Words

Text document search engines fifty years ago were already using word analysis to reduce the size of their keyword indexes. This was to manage the many variations a search term might take: MYSTERY versus MYSTERIES as well as MYSTIFY, MYSTICISM, and MYSTERIOUS. Since these all revolve around a common concept, many system builders opted to reduce them all to the single term MYSTERY in a search index. This is also helpful for maximizing the number of relevant documents retrieved for a query.

Consequently, many kinds of rule- and table-driven word stemming emerged. It can on the whole be a rather crude instrument for text processing, but might still be helpful for language analysis in general if we can do it reliably. For English at least, it turns out to be fairly straightforward if we can work long enough at the refinement of stemming. This has resulted in two quite separate PyElly tools for analyzing the structure of words as well as dealing with unknown terms.

---

### 9.2.1 Inflectional Stemming

An inflection is a change in the form of a word reflecting its grammatical use in a sentence. Indo-European languages, which include English, tend to be highly inflected; and in cases like Russian, the form of most words can vary greatly to indicate person, number, tense, aspect, mood, and case. Modern English, however, has kept only a few of the inflections of Old English, and so it has been easier to formulate rules to characterize how a particular word can vary.

PyElly inflectional stemming currently recognizes only five endings for English: -S, -ED, -ING, -N, and -T. These each have their own associated stemming logic and also share

additional logic related to recovering the uninflected form of a word. All that logic is based on American English spelling rules and special cases. PyElly coordinates its execution through the module `inflectionStemmerEN.py`.

If an unknown word ends in -S, -ED, -ING, -N, or -T, PyElly will apply the logic for the ending to see whether it is an inflection and, if so, what the uninflected word should be. Though such logic is necessarily incomplete, it has been refined by forty years of use in various systems and is generally accurate for American spellings of most English words. For example,

```
winnings ==> win -ing -s
placed   ==> place -ed
judging  ==> judge -ing
cities   ==> city -s
bring    ==> bring
sworn    ==> swear -n
meant    ==> mean -t
```

PyElly stemming will automatically prepend a hyphen (-) on any split off word ending so that it can be recognized. The original word in the PyElly input stream is then replaced by the uninflected word followed by the removed endings as shown. All the endings will be taken as separate tokens in PyElly parsing.

In some applications, you may just want to ignore the removed word endings, but these can be quite valuable for figuring out unknown words. The -ED, -ING, -N, and -T endings indicate a verb, and you can provide grammar rules to exploit that syntactic information. For example,

```
D:-ED <- ED
—
D:-T   <- ED
—
D:-N   <- ED
—
G:VERB[|ED]->UNKN ED
—
```

To use English inflectional stemming in PyElly, setting the `language` variable in the `ellyConfiguration.py` file to `EN`. To override such stemming just for a particular word, define that word in a PyElly `D:` grammar rule or a vocabulary table entry so that it will no longer be unknown.

The logic for an ending `x` is defined by a text file `x.sl` loaded by PyElly at runtime. You can also define your own inflectional stemming logic by editing the current `*.sl` files or by writing new ones. The current files for English are `Stbl.sl`, `EDtbl.sl`, `INGtbl.sl`, `Ntbl.sl`, `Ttbl.sl`, `rest-tbl.sl`, `spec-tbl.sl`, and `undb-tbl.sl`. To do inflectional stemming for a new language `ZZ`, you will have to write the `*.sl` files and a `inflectionStemmerZZ.py`. Use `inflectionStemmerEN.py` as a model here.

Here is a segment of actual logic from `Stbl.sl`, which tells PyElly what to check in identifying a -S inflectional ending when it is preceded by an IE. The literal strings for comparison in the logic below have their characters in reverse order because PyElly will be matching from the end of a word towards its start.

```
IF ei
  IF tros {SU}
  IF koo {SU}
  IF vo {SU}
  IF rola {SU}
  IF ppuy {SU}
  IF re
    IF s
      IF im {SU 2 y}
      END {FA}
    IF to {SU}
    END {SU 2 y}
  IF t
    IS iu {SU 2 y}
    LEN = 6 {SU}
    END
  END {SU 2 y}
```

This approximately translates to

```
if you see an IE at the current character position, back up and
  if you then see SORT, succeed.
  if you then see OOK, succeed.
  if you then see OV, succeed.
  if you then see ALOR, succeed.
  if you then see YUPP, succeed.
  if you then see ER, back up and
    if you then see S, back up and
      if you then see MI, succeed, but drop the word's last two letters and add Y.
      otherwise fail.
    if you then see OT, then succeed.
    otherwise succeed, but drop the word's last two letters and add Y.
  if you then see T, then back up and
    if you then see a I or a U, then succeed, but drop the word's last two letters and add Y.
    if the word's length is 6 characters, then succeed.
    otherwise succeed, but drop the word's last two letters and add Y.
```

This stemming logic is equivalent to a finite state automaton (FSA). Its operation should be fairly transparent, although the total number of different rules for English inflections has grown to be quite extensive. You may nevertheless eventually run into a case that is handled incorrectly and will want to add to the rules. Make sure, however, to test out every change so that you can avoid making everything worse.

---

## 9.2.2 Morphology

Morphology in general is about how words are put together, including processes like BLACK + BIRD ==> BLACKBIRD, EMBODY + -MENT => EMBODIMENT, and KOREAN + POP ==> K-POP. PyElly morphological analysis is currently limited to that involving the addition of prefixes or suffixes to a root, which is not necessarily a word.

The morphology component of PyElly started out as a simple FSA stemmer that served just to remove common endings from English words, including -S, -ED, and -ING. It has now evolved to focus on non-inflectional endings and to output the actual affixes removed as well as the final root form.

Earlier above, we saw “unrealizable” broken down into UN+, +REAL, -IZE, and -ABLE. True morphological analysis here would also tell us that the -IZE suffix changes the word REAL into a verb, the -ABLE suffix changes the verb REALIZE into an adjective again, and the prefix UN+ negates the sense of the adjective REALIZABLE. This is what PyElly can now do, which is useful in figuring out the syntactic type of unknown words.

For an application A, PyElly will work with prefixes and suffixes through two language rule tables defined by files `A.ptl.elly` and `A.stl.elly`, respectively. These are akin to the grammar, macro substitution, and word pattern tables already described. We have two separate files here because suffixes tend to be more significant for analyses than prefixes, and it is common to do nothing at all with prefixes.

PyElly morphological analysis will be applied only words that otherwise would be assigned the UNKN syntactic type after all other lookup and pattern matching has been done. The result will be similar to what we see with inflectional stemming; and to take advantage of them, you will also have to add the grammar rules to recognize prefixes and suffixes and incorporate them into an overall analysis of an input sentence.

---

### 9.2.2.1 Word Endings

PyElly suffix analysis will be done after any removal of inflectional endings. For application A, the `A.stl.elly` file guiding this will contain a series of patterns and actions like the following:

```
abular 2 2 le.
dacy 1 2 te. 1
entry 1 4 .
gual 2 3 . 0a
ilitation 2 6 &,
ion 2 0 .
lenger 2 5 . 0e
oarsen 1 5 .
```

```
piracy 1 4 te. 1
santry 1 4
tention 1 3 d.
uriate 2 2 y.
worship 0 0 .
|carriage 0 0 .
|safer 1 5 . 0e
```

Each line of a \*.stl.elly file defines a single pattern and actions upon matching. Its format is as follows from left to right:

- A word ending to look for. This does not have to correspond exactly to an actual morphological suffix; the actions associated with an ending will define that suffix. The vertical bar (|) at the start of a pattern string matches the start of a word.
- A single digit specifying a contextual condition for an ending to match: 0= always reject this match, 1= no conditions, 2= the ending must be preceded by a consonant, and 3= the ending must be preceded by a consonant or U.
- A number specifying how many of the characters of the matched characters to keep as a part of a word after removal of a morphological suffix. A starting vertical bar (|) in a listed ending will count as one character here.
- A string specifying what letters to add to a word after removal of a morphological suffix. An & in this string is conditional addition of e in English words, applying a method defined in English inflectional stemming.
- A period (.) indicates that no further morphological analysis be applied to the result of matching a suffix rule and carrying out the associated actions; a comma (,) here means to continue morphological analysis recursively.
- A number indicating how many of the starting characters of the unkept part of a matching ending to drop to get a morphological suffix to be reported in an analysis.
- A string specifying what letters to add to the front of the reduced unkept part of a matching ending in order to make a complete morphological suffix.

In applying such pattern rules to analyze a word, PyElly will always take the longest match. For example, if the end of a word matches the LINGER pattern above, then PyElly will ignore the shorter matches of a ENGER pattern or a GER pattern.

In the LINGER rule above, PyElly will accept a match at the end of word only if preceded by a consonant in the word. On a match, the rule specifies to keep 5 of the matched characters in the resulting root word. From the rest of a matched ending, PyElly will drop no characters, but add an E in front to get the actual suffix removed.

So the word CHALLENGER will be analyzed as follows according to the suffix patterns above:



<b>CHAL Lenger</b>	(split off matched ending and check preceding letter)
<b>CHALLENGE R</b>	(move five characters of matched ending to resulting word)
<b>CHALLENGE ER</b>	(add E to remaining matched ending to get actual suffix -ER)

The period (.) in the action for Lenger specifies no further morphological analysis. With a comma (,), PyElly would continue, possibly producing a sequence of different suffixes by reapplying its rules to the word resulting from preceding analyses. This can continue indefinitely, with the only restriction being that PyElly will stop trying to remove endings when a word is shorter than three letters.

To handle the stripped off morphological suffixes in a grammar, you should define rules like

```
D:-ion <- SUFFIX[:NOUN]
```

—

and then add G: grammar rules for dealing with these syntactic types as in the case of inflections. For example,

```
G:NOUN->UNKN SUFFIX[:NOUN]
```

—

A full grammar would of course have to be ready to deal with many different morphological suffixes.

The PyElly file `default.stl.elly` is a comprehensive compilation of English word endings evolving over the past fifty years and covering most of the non-foreign irregular forms listed in WordNet exception files. If there is more than possible analysis, PyElly will do nothing; this is so that RENT is not reduced to REND. If you actually want to make a decision here, then you must supply your own grammar rule to do it.

The `default.stl.elly` file also includes transformations of English irregular inflectional forms, which actually involve no suffix removal. For example, DUG becomes DIG -ED. This cannot be handled by PyElly inflectional stemming logic.

---

### 9.2.2.2 Word Beginnings

For prefixes, PyElly works with patterns exactly as with suffixes, except that they are matched from the beginning of a word. For example

```
contra 1 0 .
hydro 1 0 .
non 2 0.
noness 1 3.
pseudo 1 0 .
quasi 1 0 .
```

```
retro 1 0 .  
tele 1 0 .  
trans 1 0 .  
under 1 0 .
```

The format for patterns and actions here is the same as for word endings. As with endings, PyElly will take the action for the longest pattern matched at the beginning of a word being analyzed.

Prefixes will be matched after suffixes and inflections have been removed. Removing a prefix must leave at least three characters in the remaining word. Actions associated with the match of a prefix will typically be much simpler than those for suffixes, and rules for prefixes will tend to be as simple as those in the example above.

PyElly removal of prefixes will be slightly different from for suffixes. With suffixes, the word **STANDING** becomes analyzed as **STAND -ING**, but with the prefix rules above, **UNDERSTAND** would become **UNDER+ +STAND**. Note that a trailing + is used to mark a removed prefix instead of a leading – for suffixes.

In the overall scheme of PyElly processing of an unknown word, inflections are checked first, then suffixes, and finally prefixes. If there is any overlap between the suffixes and the prefixes here, then inflections and suffixes takes priority.

For example, **NONFUNCTIONING** becomes **NON+ +FUNCT -ION -ING** with the morphology rules above. A grammar would then have to stitch these parts back together in an analysis.

For prefixes here, you will need a dictionary rule like

```
D:non+ <- PREFIX[+NEG]
```

—

and you should by now know how to supply the required grammar rules yourself.

## 9.3 Entity Extraction

In computational linguistics, an entity is some phrase in text that stands for something specific that we can talk about. This is often a name like George R. R. Martin or North Carolina or a title like POTUS or the Bambino; but it also can be insubstantial as Flight VX 84, 888-CAR-TALK, 2.718281828, NASDAQ APPL, or orotidine 5'-phosphate.

The main problem with entities is that we are likely to have almost none of them in a predefined vocabulary. People seem to handle them in stride while reading text, however, even when they are unsure what a given entity means exactly. This is in fact the purpose of much text that we read: to inform us about something we might be unfamiliar with. A fully competent natural language system must be able to function in this kind of situation.

At the beginning of the 21st Century, systems for automatic entity extraction from text were all the rage for a short while. Various commercial products with impressive capabilities came on the market, but unfortunately, just identifying entities is insufficient to build a compelling application, and so entity extraction systems mostly fell by the wayside in the commercial marketplace. In a tool like PyElly, however, some builtin entity extraction support can be quite valuable.

---

### 9.3.1 Numbers

PyElly no longer has a predefined `NUM` syntactic type. The PyElly predecessor written in C did have compiled code for number recognition, but this covered only a few possible formats and was dropped later in Jelly and PyElly for a more flexible solution. If you want PyElly to recognize literal numbers in text input, you must make use of special patterns in files `*.p.elly` as described in Section 4.2.

PyElly, however, also has gone further here. It also has some builtin capabilities for automatic normalizations of number references so that you need fewer patterns to recognize them. In particular,

- Automatic stripping out of commas in numbers as an alternative to doing this with special pattern matching:

1,000,000 ==> 1000000.

- Automatic mapping of spelled out numbers to a numerical form:

one hundred forty-third ==> 143rd

fifteen hundred and eight ==> 1508

Here you still need patterns to recognize the rewritten numbers so that PyElly can process them. You can disable all such number rewriting by setting the variable `ellyConfiguration.rewriteNumbers` to `False`.

---

### 9.3.2 Dates and Times

Dates and Times could be handled as PyElly patterns, but their forms can vary so much that this would take an extremely complicated finite-state automaton. For example, here are just two of many possible kinds of dates:

the Fourth of July, 1776  
2001/9/11

To recognize such entities, the PyElly module `extractionProcedure.py` defines some date and time extraction methods written in Python that can be called automatically when processing input text.

To make such methods available to PyElly, they just have to be listed in the `ellyConfiguration.py` module. Here is the actual Python code to do so:

```
import extractionProcedure

extractors = [ # list out extraction procedures
    [ extractionProcedure.date , 'date' ] ,
    [ extractionProcedure.time , 'time' ]
]
```

You can disable date or time extraction by just removing its method name from the `extractors` list. The second element in each listed entry is a string syntax specification, which generally includes a syntactic type plus syntactic features to assign to a successfully extracted entity. The names of syntactic types and features here will have to be coordinated with other PyElly grammar rules.

The date and time methods above are part of the standard PyElly distribution. These will do some normalization of text before trying to recognize dates and times. Dates will be rewritten in the form

`mm/dd/yyyyXX`

For example, 09/11/2001AD. Times will be converted to a 24-hour notation

`hh:mm:ssZZZ`

For example, 15:22:17EST. If date or time extraction is turned on, then your grammar rules should expect to see these forms when a generative semantic procedure executes an `OBTAIN` command. The `XX` epoch indicator in a date and the `ZZZ` zone indicator in a time may be omitted in PyElly input.

---

### 9.3.3 Other Entities

You can write your own entity extraction methods in Python and put them into the `extractors` list for PyElly. This should be done as follows:

1. The name can be anything legal in Python for method names.
2. The method should be defined at the level of a module, outside of any Python class. This should be in a separate Python source file, which can then be imported into `ellyConfiguration.py`.
3. The method takes a single argument, a list of individual Unicode characters taken from the current text being analyzed. PyElly will prepare that list. The method may alter the list as a side effect, but you will have to be careful in how you do this if you want the changes to persist after returning from the method. That is because Python always passes arguments to a method by value.

4. The method returns the count of characters found for an entity or 0 if nothing is found. The count will always be from the start of an input list after any rewriting. If no entity is at the current position in input text, return 0.
5. PyElly will always apply entity extraction methods in the order that they appear in the extractors list. Note that any rewriting of input by a method will affect what a subsequent method will see. All extractor methods will be tried.
6. An extraction method will usually do additional checks beyond simple pattern matching. Otherwise, you may as well just use PyElly finite-state automata described in Section 4.2.
7. Install a new method by editing the extractors list in the PyElly file `ellyConfiguration.py`. You will have to import the module containing your method.

The module `extractionProcedure.py` defines the method `stateZIP`, which looks for a U.S. state name followed by a five- or nine-digit postal ZipCode. This will give you a model for writing your own extraction methods; it is currently not installed.

## 9.4 PyElly Vocabulary Tables

PyElly maintains large vocabulary tables in external files created with the free Berkeley Database package, which is separate from Python. To use vocabulary tables, you must download and install Berkeley Database on your computing platform along with the third-party `bsddb3` Python module that provides an interface between Python and Berkeley Database access functions written in C.

Installing Berkeley Database and `bsddb3` can be fairly straightforward, but depends on the target platform and so will not be described here. You may want to call on some expert technical support. Appendix C shows how this was done on Mac OS X.

You can run PyElly without vocabulary tables, but these can make life easier for you even when working with only a few hundred different terms. They provide the most convenient way to handle multi-word terms and terms including punctuation. They also can be reused with different grammar tables and generally will be more compact and easier to compile than D: rules of a grammar. Without them, PyElly will be limited to very simple applications.

PyElly vocabulary table entries in a `*.v.elly` definition file in which each entry can be only a single text line and can have only extremely limited semantics. This is in large part so that one may generate large volumes of entries automatically through scripts. For example, the PyElly distribution file `default.v.elly` has 155,229 entries generated with bash shell scripts from WordNet 3.0 data files.

Each PyElly vocabulary table entry must take one of the following formats:

TERM : SYNTAX

TERM : SYNTAX =TRANSLATION

TERM : SYNTAX x=Tx, y=Ty, z=Tz

TERM : SYNTAX (procedure)

TERM : SYNTAX SEMANTIC-FEATURES PLAUSIBILITY

TERM : SYNTAX SEMANTIC-FEATURES PLAUSIBILITY =TRANSLATION

TERM : SYNTAX SEMANTIC-FEATURES PLAUSIBILITY x=Tx, y=Ty, z=Tz

TERM : SYNTAX SEMANTIC-FEATURES PLAUSIBILITY (procedure)

The TERM : SYNTAX part is mandatory for vocabulary entry. A TERM can be

Lady Gaga

Lili St. Cyr

Larry O'Doule

“The Robe”

ribulosebiphosphatecarboxylaseoxygenase

The ‘:’ is required to let PyElly know when a term ends; no wildcards are allowed in a TERM, and it must start with a letter, digit, or the character ‘.’ or ‘\’. SYNTAX is just the usual PyElly specification of syntactic type plus optional syntactic features.

SEMANTIC-FEATURES are the bracketed semantic features for cognitive semantics (see Subsection 10.2); it can be “0” or “-” if no features are set. PLAUSIBILITY is an integer value for scoring a phrase formed from a vocabulary entry; this value may include an attached semantic concept name separated by a “/” (see Subsection 10.3). Both SEMANTIC-FEATURES and PLAUSIBILITY may be omitted, but if either is present, then the other must be also.

The final translation part of a vocabulary entry is optional and can take one of the forms shown above. If the translation is omitted, then the generative semantic procedure for the entry will be just the operation OBTAIN. This is equivalent to no translation at all.

An explicit TRANSLATION is a literal string to be used in rewriting a vocabulary entry; the ‘=’ is mandatory here. The x=Tx, y=Ty, z=Tz alternate form is a generalization of the simpler TRANSLATION; it maps to the generative semantic operation

PICK lang (x=Tx#y=Ty#z=Tz#)

It is possible here that one of the `x` or `y` or `z` in the translation options of a vocabulary entry can be the null string. In this case, the `PICK` operation will treat the corresponding translation as the default to be taken when the value of the `lang` PyElly local variable matches none of the other specified options.

A (procedure) in parentheses is a call to a generative semantic subprocedure defined elsewhere in a `*.g.elly` grammar rule file.

Here are some full examples of possible vocabulary table entries in a `*.v.elly` file:

```
Lady Gaga : noun [^celeb] =Stefani Joanne Angelina Germanotta
Lili St. Cyr : noun[:name] [^celeb] 0
horse : noun FR=cheval, ES=caballo, CN=馬, RU=лошадь
twerk : verb[|intrans] [^sexy] (xxxx)
```

All references to syntactic types, syntactic and semantic features, and procedures will be stored in a vocabulary table as an encoded numerical form according to a symbol table associated with a PyElly grammar.

Syntactic features must be immediately after a syntactic category name with no space in between. Otherwise, PyElly will be unable to differentiate between syntactic and semantic features in a `*.v.elly` file. Individual syntactic or semantic features inside of brackets may be preceded by a space, however.

Unlike the dictionary definitions of words in a grammar, there are no permanent rules associated with the terms in a vocabulary table. When a term is found by lookup, PyElly automatically generates a temporary internal dictionary rule to define the term. This rule will persist only for the duration of the current sentence analysis.

A term may have multiple vocabulary table entries; for example,

```
bank : noun [^institution]
bank : noun [^geology]
bank : verb [|intrans]
```

If the word `BANK` shows up in input text, then all of these entries will be tried out in possible PyElly analyses.

Often, a vocabulary table may have overlapping entries like

```
manchester : noun [^city]
manchester united : noun [^pro,soccer,team]
```

PyElly will always take the longest matching entry in the analysis of an input sentence and ignore any shorter matches.

For a given application `A`, the PyElly `ellyMain` module will look for a vocabulary table definition in the text file `A.v.elly`. If this is missing, the file `default.v.elly` is taken; this is a listing of most of the nouns, verbs, adjectives, and adverbs in WordNet

3.0. Define an empty `A.v.elly` file if you do not want this huge default vocabulary, which will take a long time to read in and compile. PyElly will always save a compiled vocabulary data base for an application `A` in the file `A.vocabulary.elly.bin`. If you change `A.v.elly`, PyElly will automatically recompile any `A.vocabulary.elly.bin` at startup. Recompile will also happen if `A.v.elly` has changed. Otherwise, PyElly will just read in the last saved `A.vocabulary.elly.bin`.

Note that the `A.vocabulary.elly.bin` file created by PyElly must always be paired only with the `A.rules.elly.bin` file it was created with. This is because syntactic types and features are encoded as numbers in `*.elly.bin` files, which may be inconsistent when they are created at different times. If you want to reuse language rules, always start from the `*.*.elly` files. If PyElly has to recompile `A.rules.elly.bin` at startup, then it will automatically recompile `A.vocabulary.elly.bin`.



## 10. Logic for PyElly Cognitive Semantics

The generative semantic part of a grammar rule tells PyElly how to translate its input into output, while the cognitive semantic parts of different grammar rules help in evaluating the plausibility of analyses. Generative semantics is always involved in producing final PyElly output; cognitive semantics comes into play only when we have to choose between alternate analyses of the same input, but it is always run for each new phrase node created in a PyElly analysis to get a plausibility score anyway.

With large grammars, we cannot expect that every input sentence will always break down in only one way into constituents according to the rules of that grammar. In most languages, for example, a particular word occurrence could be assigned multiple parts of speech, and each possibility here can result in a different syntactic analysis for an input sentence. Those alternate analyses can potentially lead to conflicting interpretations of a sentence that PyElly eventually has to resolve somehow.

PyElly takes a conservative approach in handling such ambiguous situations. Multiple interpretations might exist at lower levels of a parse tree, but some could end up not fitting into any final analysis for an entire sentence. In this way, an ambiguity might resolve itself once we can see the bigger context. So PyElly holds off on making any decision until it has more information.

PyElly will disregard ambiguities until its analysis produces two or more phrase nodes of the same syntactic type with the same syntactic features over the same segment of an input sentence. At that point, PyElly examines the plausibility score for each alternate phrase node according to its cognitive semantics and then choose just one alternate to keep in continuing to build a full sentence analysis.

Note that some kinds of ambiguity can slip completely past PyElly here. For example, “I love rock” could be about music or landscaping. If your grammar rules fail to give you two different syntactic analyses with co-extensive phrase nodes of the same type and same features, however, PyElly will be unaware of alternative interpretations here. PyElly knows about language only what you can tell it explicitly.

Here is a simple example. Consider the following grammar rules:


```

g:sent[:*r]->x
—
g:sent[:*r]->y
—
d:www<-x[:f]
—
  (xgen)
—
d:www<-y[:g]
—
  (ygen)
—

```

where  $w_{www}$  is a dictionary word that can be associated with two different syntactic types  $x$  and  $y$ . A sentence consisting only of the word  $w_{www}$  will therefore be ambiguous at the lowest level of analysis because we do not know whether the proper generative semantics for the sentence should call  $(x_{gen})$  or  $(y_{gen})$  as a subprocedure.


We have two possible analyses here, given the rules for inheritance of syntactic features through the predefined  $*R$  syntactic feature described in Subsection 8.1:



PyElly cannot resolve any ambiguities here, however, because none of the constituents in the two alternate analyses of the sentence “ $w_{www}$ ” have the same syntactic type and the same syntactic features. Now if we replace the first two rules above with

```
g:sent->x
—
g:sent->y
—
```

our possible analyses instead become:



There is now a match for interpretations at the  $SENT$  level, and PyElly can now decide which one of them to choose for its final interpretation of the sentence “ $w_{www}$ ”. You can supply the cognitive semantics here to favor one or the other.

Remember that PyElly should not select between interpretations at the  $X[:f]$  and  $Y[:g]$  level because they are not syntactically interchangeable. An analysis could still ultimately fail with one or the other, however, and so we really need to carry along both options so as not to preclude a successful analysis later.

This section of the User's Manual will tell you how to score alternative analyses when your grammar rules do allow for ambiguities. You can still write language descriptions without such explicit scoring, but it is built into the PyElly parsing algorithm, and this is just a way to give you more control over how PyElly runs.

The PyElly plausibility score for an analyzed constituent of a sentence will be a simple integer value. A zero will be neutral, and positive and negative will indicate degrees of

plausibility and implausibility, respectively. The score for a particular phrase node generally will add up the scores of its immediate phrase subconstituents plus a contribution from the cognitive semantics of the grammatical rule combining those subconstituents into one resulting phrase.

For example, suppose we have a constituent described by a grammar rule  $A \rightarrow X \ Y$ . We first get a plausibility score for subconstituent  $X$  and another score for subconstituent  $Y$ ; with PyElly, these would be computed already. Then we run the cognitive semantic logic for the grammar rule  $A \rightarrow X \ Y$ , which produces a value to add to the plausibility scores for  $X$  and  $Y$  to get and save an overall plausibility for our phrase of type  $A$ .

With competing analyses, if only one has the top score, PyElly chooses it and is done. If more than one has the highest score, then PyElly will arbitrarily pick one of them. PyElly will keep track of which grammar rules are involved here, however, and will make a different choice the next time an ambiguity arises with these rules.

## 10.1 The Form of Cognitive Semantic Clauses

PyElly currently provides three different ways to define how a grammar rule will contribute to a plausibility score: by a fixed value assigned to the grammar rule in a language definition, by logical rules relating the semantic features associated with the constituents to be combined by the grammar rule, and by measuring the semantic distance between the concepts associated with each constituent.

In the input `*.g.elly` file defining grammar rules for a PyElly language description, the cognitive semantic logic will consist of a series of single-line clauses coming after the `G:` or a `D:` line introducing each rule and ending at the first `_` or `__` line (see Section 4). Each clause will be a line containing the character sequence `'>>'`, possibly with text before and after. For example, the following rule with no explicit generative semantics has three cognitive semantic clauses:

```
G:NP->ADJ NOUN
  L[^EXTENS] R[^ABSTRACT]>>*R-
  R[^GENERIC] >>*L+
  >>*R
_
```

A `'>>'` divides a clause into two major parts. The left part specifies conditions on the immediate subconstituents of a phrase structure in order for the clause to apply. The right part specifies actions to take if the left side is satisfied. The `'>>'` is mandatory in a cognitive semantic clause.

When evaluating the plausibility of a given rule of grammar, each of its cognitive semantic clauses will be tried in order until one is found to apply. None may apply, in which case a zero contribution is assumed. Having no conditions on the left side of a clause is equivalent to an always-true condition, and such a clause will always make any following it irrelevant.

The actual form of a clause will depend on which of the three kinds of plausibility contribution is being specified in the clause. You may freely mix the different kinds within the same clause or the same set of clauses, but remember that ordering does matter here; the first clause to match will always define the contribution of a grammar rule for an overall plausibility score.

A grammar rule may have cognitive semantic clauses even it has no explicit generative semantic procedure. In this case, the listing of clauses will be terminated by a double underscore (\_\_) line without a preceding single underscore (\_).

## 10.2 Cognitive Semantic Elements

PyElly cognitive semantics shows up mainly in the grammar rule logic for evaluating the plausibility of phrase nodes in an analysis, but appears in various other places as well. There are three different elements involved: fixed scoring, semantic features, and semantic concepts.

---

### 10.2.1 Fixed Scoring

The simplest and most common kind of cognitive semantic clause will assign a fixed positive or negative score unconditionally to a grammar rule in order to favor or disfavor a phrase analysis based on the rule. Such clauses may take one of the following forms:

```
>>-
>>+
>>+++
>>-----
>>+5
>>-20
```

The initial + or - signs are mandatory in the scoring. A string of n +'s or -'s is equivalent to +n or -n. Here is an example of use:

```
G:NP->ADJ UNKNOWN
>>-      # cognitive semantics disfavoring this rule by -1

- RIGHT  # generative semantics
LEFT    #

—
```

If no cognitive semantic clauses are specified for a grammar rule, this is equivalent to

```
>>+0
```

a special case of fixed scoring. Note that the “+” is necessary here if you actually want to be explicit here about a zero score.

---

## 10.2.2 Semantic Features

Semantic features are similar to syntactic features as defined above in Section 8, but play no role in distinguishing between different grammar rules. They are specified in the same bracketed format as syntactic features; for example:

```
[ &ANIMATE, MOBILE ]
```

The & is the feature set identifier, and ANIMATE and MOBILE are two specific features. Semantic features will have completely separate lookup tables from syntactic features. In particular, a syntactic feature set and a semantic feature set can have the same set identifier without any conflict, but always make the identifiers different just for clarity.

As with syntactic features, you may have up to 16 semantic feature names, with the rules for legal names being the same. Unlike syntactic features, however, they have no predefined feature names like \*LEFT or \*RIGHT.

---

### 10.2.2.1 Semantic Features in Cognitive Semantic Clauses

A cognitive semantic clause for a 2-branch splitting G: grammar rule will have the following general form

$$L[oLF1, \dots, LFn] \ R[oRF1, \dots, RFn] >> x[oF1, \dots, Fn] \#$$

The symbol “o” is a feature set identifier; “x” may be either \*L or \*R, and the “#” is a fixed scoring as in Subsection 10.1 above; for example, +++ or -3.

The prefixes L and R on the left part of a clause specify the constituent substructures to be tested, respectively left and right. Either can be omitted, but you probably want to specify at least one. Otherwise, you may as well use fixed scoring.

The “x” prefix on the right is optional for specifying inheritance of features. An \*L means to copy the semantic features of the left subconstituent into the current phrase; \*R means to copy the right. You cannot have both, but a missing “x” means no inheritance at all. Any explicit semantic feature set in the right part of the clause will indicate any additional features to turn on for a phrase node.

A cognitive semantic clause for a 1-branch extending G: grammar rule will have the following general form in its semantic features:

$$L[oLF1, \dots, LFn] >> x[oF1, \dots, Fn] \#$$

A D: grammar rule defines a phrase without any constituent substructures. So a semantic features in a clause must take the form

```
>>[oF1, . . . , Fn] #
```

That is, you can set semantic features for a `D:` rule, but not test or inherit any.

For both splitting and extending grammar rules, any of the left side of a cognitive semantic clause can be omitted. If all are omitted here, then a clause applies unconditionally.

---

### 10.2.2.2 Semantic Features in Generative Semantics

PyElly also allows a generative semantic procedure to look at the semantic features for the phrase node to which it is attached. This is done in a special form of the `IF` command where the testing of a local variable is replaced by the testing of semantic features. For example,

```
IF  [&F1, -F2, F4]
    (DO-SOMETHING)
END
```

The testing here is exactly the same as the testing of syntactic features in determining the applicability of a 1- or 2-branch grammar rule in PyElly parsing.

---

### 10.2.3 Semantic Concepts

PyElly is currently being used experimentally in applying conceptual information from WordNet to infer the intended sense of ambiguous words in English text. This has been incorporated into a new cognitive semantic option in PyElly.

PyElly allows you to establish a set of concepts each identified by a unique alphanumeric string and related to one other by a conceptual hierarchy defined in a language description for an application. This can be done any way that you want, but WordNet provides a good starting point here since it contains over two hundred thousand different synonym sets (or synsets) as potential concepts to work with.

(WordNet is produced manually by professional lexicographers affiliated with the Cognitive Science Laboratory at Princeton University and is an evolving linguistic resource now at version 3.1. [George A. Miller (1995). WordNet: A Lexical Database for English. Communications of the ACM Vol. 38, No. 11: 39-41.] This notice is required by the WordNet license.)

In WordNet, each possible dictionary sense of a term will be represented as a set of synonyms (synset). This can be uniquely identifiable as an offset into one of four data files associated with the main parts of speech—`data.noun`, `data.verb`, `data.adj`, and `data.adv`.

For disambiguation experiments in PyElly, trying to work with all the synsets of WordNet 3.1 is too cumbersome. So we instead have been focusing on concepts from a small subset of WordNet synsets related to interesting kinds of ambiguity in English. We can identify each such concept as an 8-digit decimal string combining the unique WordNet offset for its corresponding synset plus a single appended letter to indicate its part of speech. For example,

```
13903468n : (=STAR) a plane figure with 5 or more points; often used as an emblem  
01218092a : (=LOW) used of sounds and voices; low in pitch or frequency
```

The standard WordNet coding for part of speech is n = noun, v = verb, a = adjective, r = adverb.

For any set of such concepts, we can then map selected semantic relations for them from WordNet into a simple PyElly conceptual hierarchy structure, which will be laid out in a PyElly language definition file `A.h.elly`. The current `disambig` example application in the PyElly package has a hierarchy with over 800 such related concepts, all taken from WordNet 3.1.

You can of course also define your own hierarchy of concepts with their special hierarchy of semantic relations. The only restriction here is that each concept name must be an alphanumeric string like `AA0123bcdefoo`. Upper and lower case will be ignored in letters. Such semantic concepts can be explicitly employed by cognitive semantic clauses on their left side and can be explicitly employed in one way and implicitly employed in two ways on the right side.

---

### 10.2.3.1 Concepts in Cognitive Semantic Logic

Semantic concepts serve to provide another contribution when computing plausibility scores to choose between alternate interpretations in case of ambiguity. This will happen in the cognitive semantic logic associated with each syntax rule, and the concepts will have different roles when they are on the left and on the right sides of a cognitive semantic clause.

---

#### 10.2.3.1.1 Concepts in the Left Side of Cognitive Semantic Clauses

The left half of a clause is for testing its applicability to a particular phrase, and PyElly allows the semantic concepts associated with its subconstituents to be checked out. The syntax here is similar to how you test semantic features of subconstituents, except you will use parentheses ( ) to enclose a concept name instead of the [ ] around semantic features. Here is an example of a concept check:

```
L(01218092a) R(13903468n) >>+
```

This checks whether the left subconstituent of a phrase has a concept on a path down from concept 01218092a in a conceptual hierarchy and whether the right subconstituent has a concept on a path down from concept 13903468n. The ordering of testing here does not matter, and you may omit either the `L` or the `R` test or both.

You can mix concept testing with semantic feature testing in the conditional part of a cognitive semantic clause. For example,

```
L(01218092a) L[^PERSON] >>++
```

You may also specify more than one concept per test. For example,

```
L(00033319n,08586507n) >>+
```

Here, PyElly will check for either `L(00033319n)` or `L(08586507n)`.

You of course can define more self-descriptive concept names for your own application. You are not limited to WordNet 3.1 synset ID's.

---

#### 10.2.3.1.2 Concepts in the Right Side of Cognitive Semantic Clauses

A single concept can be explicitly appended on the right side of a clause with a separating space. For example,

```
>>+ CONCEPT
```

This must always come after a plausibility scoring expression. If you want a neutral scoring here, you must specify it explicitly here as

```
>>+0 CONCEPT
```

Normally, this kind of concept reference will be useful only for the cognitive semantics of `D`: dictionary rules of a grammar, but nothing keeps you from trying it out in `G`: rules as well.

Concepts can also show up implicitly of the right side of a clause. When a subconstituent of a phrase has an associated concept, the `*L` or `*R` inheritance actions specified by a clause will apply to concepts as well. So, a clause like

```
>> *L++
```

will cause not only the inheritance of semantic features from a left subconstituent, but also the inheritance of any semantic concept from that left subconstituent. That is also true for `*R` with a right subconstituent.

To use semantic concepts on the right side of a clause, you generally must use the `*L` or `*R` mechanism even if you have no semantic features defined. This must be done to pass



concepts in a parse tree for later checking. Note that you cannot have both  $*_L$  and  $*_R$  in a cognitive semantic clause.

Semantic concepts also implicitly come into play in two ways when PyElly is computing a plausibility score for a phrase:

1. When a subconstituent has a semantic concept specified, PyElly will check whether it is on a downward path from a concept previously seen in the current or an earlier sentence. PyElly will maintain a record of such previous concepts to check against. If such a path is found, the plausibility score of a phrase will be incremented by one. If a phrase has one subconstituent, the total increment possible here is 0 or 1; if the phrase has two, the total increment could be 0, 1, or 2.
2. If a phrase has two subconstituents with semantic concepts, PyElly will compute a semantic distance between their two concepts in our inverted tree by following the upward paths for each concept until they intersect. The distance here will be the number of levels from the top of the tree to the point of intersection. If the intersection is at the very top, then the distance will be zero. The lower the intersection in the tree, the higher the semantic relatedness. This distance will be added to the plausibility score of a phrase containing the two subconstituents.

If no semantic concepts are specified in the subconstituents of a phrase, then a semantic plausibility score will be computed exactly as before.

---

### 10.2.3.2 Semantic Concepts in Language Definition Files

To use semantic concepts, you must define them in a PyElly language definition. For an application A, this must happen in the files `A.h.elly`, `A.g.elly`, or `A.v.elly`. They can be omitted entirely if you have no interest in them.

---

#### 10.2.3.2.1 Conceptual Hierarchy Definition (`A.h.elly`)

This specifies all the concepts in a language definition and their semantic relationships. You can define everything arbitrarily, but to ensure consistency, start from some existing language database like WordNet. Here are some entries from `disambig.h.elly`, a conceptual hierarchy definition file based on WordNet 3.1 concepts for a PyElly example application:

```
14831008n > 14842408n
14610438n > 14610949n
00033914n > 13597304n
05274844n > 05274710n
07311046n > 07426451n
07665463n > 07666058n
04345456n > 02818735n
03319968n > 03182015n
08639173n > 08642231n
00431125n > 00507565n
```

The “>” separates two concept names to be interpreted as a link in a conceptual hierarchy, where the left concept is the parent and the right concept is a child. In this particular definition file, each concept name is an offset in a WordNet 3.1 part of speech data file plus a single letter indicating which part of speech (n, v, a, r). Both offset and part of speech are necessary to identify any WordNet concept uniquely.

For convenience, a \*.h.elly file may also have entries of the form

```
=xxxx yyyy
=zzzz wwwww
```

These let you to define equivalences of concept names, where the right concept becomes the same as the left. For example, the entries make yyyy to be the same as xxxx and wwwww to be the same as zzzz. The left concept must occur elsewhere in the hierarchy definition, though. An equivalence can be specified anywhere in a \*.h.elly file. It specifies only a convenient alias for a concept name without defining a new concept, which can be helpful in documentation of semantic relationships.

---

#### 10.2.3.2.2 Semantic Concepts in Grammar Rules (A.g.elly)

This was already discussed briefly in Subsection 10.3.2 above. Here is an example of a grammar dictionary rule with cognitive semantics referencing semantic concepts:

```
D:xxxx <- NOUN
>>+0 CONCEPT

—
  APPEND xxxx-C
—
```

Similarly with a regular syntax rule:

```
G:X -> Y Z
>>*L+0 CONCEPT

—
  RIGHT
  SPACE
  LEFT
—
```

Note here that the \*L action will also cause any concept associated with Y to be inherited by X, but the explicit assignment of CONCEPT here will always override any such inheritance.

### 10.2.3.2.3 Semantic Concepts in Vocabulary Table Entries (A.v.elly)

For a vocabulary table entry, we extend the plausibility field in an `A.h.elly` input file to allow appending a concept name separated by a “/” (See Subsection 9.4 above). Omitting a concept name here will be equivalent to a null concept.

Here are some entries from `disambig.v.elly`, a vocabulary table definition file making use of the concepts above.

```
finances : noun[:*unique] - 0/13377127n =funds0n
monetary resource : noun[:*unique] - 0/13377127n =funds0n
cash in hand : noun[:*unique] - 0/13377127n =funds0n
pecuniary resource : noun[:*unique] - 0/13377127n =funds0n
assets : noun[:*unique] - 0/13350663n =assets0n
reaction : noun[:*unique] - 0/00860679n =reaction0n
response : noun[:*unique] - 0/00860679n =reaction0n
covering : noun[:*unique] - 0/09280855n =covering0n
natural covering : noun[:*unique] - 0/09280855n =covering0n
cover : noun[:*unique] - 0/09280855n =covering0n
```

The `*UNIQUE` syntactic feature in each entry is to disable PyElly ambiguity resolution at lower levels of sentence analysis, a requirement for the `disambig` example application. The translation provided for each entry above is the WordNet offset designation for a particular word sense plus a single letter specifying its part of speech.

You may name the concepts in your own `A.h.elly` hierarchy definitions however you wish, but with two exceptions: the name “-” will be reserved to denote a null concept explicitly in grammar rules; and the name “^” will be reserved for the top of a hierarchy to which every other concept is linked eventually. You must have “^” somewhere in a `A.h.elly` hierarchy definition file for it to be accepted by `vocabularyTable.py` as a language definition file.

## 11. Sentences and Punctuation

Formal grammars typically describe the structure of single sentences. PyElly accordingly is set up to analyze one sentence at a time through its `ellyBase` module. In typical text, however, sentences are all jumbled together, and we somehow have to divide them up properly before we can do anything with them. That task is harder than one might think; for example,

```
I met Mr. J. Smith at 10 p.m. in St. Louis.
```

This sentence contains six periods (.), but only the final one stops the sentence. It is not hard to recognize such exceptions, but this is yet more detail to take care of on top of an already complex task of language analysis.

PyElly divides text into sentences with its `ellySentenceReader` module, which employs a simple sequential punctuation-checking algorithm to detect sentence boundaries in text. While doing this, PyElly also normalizes each sentence to make subsequent processing easier. The simplicity of the algorithm will make it tend to find too many sentences, but we can help it out by providing some supporting modules with more smarts.

Currently, the PyElly `stopException` module works with a user-provided list of patterns to determine whether a particular instance of punctuation like a period (.) should actually stop a sentence. The PyElly `exoticPunctuation` module, tries to normalize various kinds of unorthodox punctuation found in informal text. This solution is imperfect, but can be modified easily. See Subsection 11.1 below for details.

The approach of PyElly here is to provide sentence recognition a notch or two better than what one can cobble together just using Python regular expressions or the standard sentence recognition methods provided in Java. If you really need more than this, then there are other resources available; for example, NLTK, which can be trained on sample data to discover its own stop exceptions. PyElly sentence recognition should be adequate for most of its potential applications, however.

PyElly sentence reading currently operates as a pipeline configured as follows:

```
raw text => ellyCharInputStream => ellySentenceReader => ellyBase
```

where `raw text` is an input stream of Unicode encoded as UTF-8 and readable line by line with the Python `readline()` method. The `ellyCharInputStream` module is a filter that removes extra white space, substitutes for Unicode characters not in the Latin-1 subset, and replaces single new line characters with spaces. The `ellyCharInputStream` and `ellySentenceReader` modules both operate at the character level and serve together to divide input text into individual sentences for PyElly processing.

A single input line could contain multiple sentences, or a single sentence may extend across multiple input lines. There is no limit on how long a single input line may be; it could be an entire paragraph terminated by a linefeed as found in many word processing files. PyElly can also read text divided into short lines terminated by linefeeds, carriage

returns, or carriage returns plus linefeeds. It currently will not splice back a hyphenated word split across two lines, however.

The `ellySentenceReader` module currently recognizes five kinds of sentence punctuation: period (.), exclamation point (!), question mark (?), colon (:), and semicolon (;). By default, these will always be taken as a sentence stop whenever followed by a whitespace character. A blank line consisting of two new line characters together will also terminate a sentence without any punctuation. The procedure here is deliberately minimal.

The `ellyMain` module, the standard command line interface for PyElly, employs `ellySentenceReader`. This can be run interactively from a keyboard, but since it expects general text input, you may have to press an extra `<RETURN>` to get PyElly to recognize the end of a sentence in order to start processing.

## 11.1 Extending Sentence Recognition

The division of text into sentences by `ellySentenceReader` can currently be modified in two ways: by the `stopException` module that recognizes special cases when sentence punctuation should not terminate a sentence and by the `exoticPunctuation` module that checks for cases where sentence punctuation can be more than a single character.

---

### 11.1.1 Stop Punctuation Exceptions

When PyElly starts up an application `A`, its `stopException` module will try to read in a file called `A.sx.elly`, or failing that, `default.sx.elly`. This file specifies various patterns for when a text character should not be treated as normal sentence punctuation.

The patterns in a `*.sx.elly` file must each be expressed in the following form:

```
l...lp|r
```

where `p` is the punctuation character for the exception, `l...l` is a sequence of character for the immediate left context of `p`, and `r` is the immediate right context character of `p`. The vertical bar (`|`) marks the start of a right context; if it is missing, the right context is assumed to be any nonalphanumeric character.

The `l` and `r` parts of a pattern may be Elly wildcards for matching. Those currently recognized in `stopException` are

- `_` matches a single whitespace character or beginning or end of text
- `@` matches a single letter
- `#` matches a single digit
- `~` matches a single nonalphanumeric character

A left context may have any of these wildcards; a right context can recognize only the whitespace wildcard (`_`). All nonwildcard characters in a pattern must be matched exactly, except for letters, which will be matched irrespective of case.

Here are some examples of actual exception patterns from `default.sx.elly`:

```
~@. | _
DR.
MR.
MRS.
U.S.S. | _
U.S. | _
```

The first pattern picks up initials, which consist of a single letter followed by a period and a space character. The other patterns match personal titles and work as you would expect them to. The file `default.sx.elly` has an extensive list of stop exceptions that might be helpful for handling typical text. You can of course supply your own list.

You should note that ordering makes a difference in the listing of patterns here. PyElly will always take the first match, which should do the right thing. You should, however, watch out for patterns where it makes a difference what character precedes the match. In the case of `DR.`, the preceding character probably does not matter; but in the case of `@.`, it will. This is why the pattern here needs to be `~@.`

---

### 11.1.2 Exotic Punctuation

This is for dealing with punctuation like `!!!` or `!?`. The capability is coded into the PyElly exoticPunctuation module, and its behavior cannot be modified except by changing the Python logic of the module. This is not hard, though.

The basic procedure here is to look for contiguous sequences of certain punctuation characters in an input stream. These are then automatically collapsed into a single character to be passed on to the `ellySentenceReader` module. The main `ellyBase` part of PyElly should therefore always see standard punctuation.

## 11.2 Predefined Punctuation Syntax

The input sentences processed by `ellyBase` will currently include all punctuation, including those recognized by `stopException` as not breaking a sentence. When PyElly breaks a sentence into parts for analysis, a single punctuation character by default will be taken as a token. PyElly will assign common English punctuation to the predefined syntactic type `PUNC` unless you provide vocabulary table rules or `D:` grammar rules or FSA pattern rules specifying otherwise.

For example, you might want to put `DR.` into your vocabulary table, perhaps as the syntactic type `TITLE`. Since this takes three characters from an input stream, including

the period, PyElly will no longer see the punctuation here. PyElly tokenization will always take longest possible match when multiple PyElly rules can apply.

The identification of punctuation as a sentence part is just the start of PyElly analysis. The grammar rules for a PyElly application will then have to describe how to fit the punctuation into the overall analysis of a sentence and how eventually to translate it. This will be entirely your responsibility; and it can get complicated.

In simple text processing applications taking only a sentence at a time, you might choose just to ignore all punctuation, but in others, punctuation occurrences in sentence will provide important clues about the boundaries of phrases in text input. In the former case, you can have a grammar rule like

```
g: UNKN->PUNC
```

---

or alternatively, define macro substitutions to make all punctuation marks disappear from input text.

When you have to work with sentence punctuation, you will need at least one grammar rule like

```
g: SENT->SENT PUNC
```

---

for handling stop punctuation terminating sentences, plus other rules for handling punctuation like “,” “(, “)”, or single and double quotation marks, which may be internal to a sentence. You will have to decide how much work you want to do here.

Elly parsing will fail if any part of a sentence cannot be put into a single coherent analysis; and punctuation handling will be a highly probable point of failure here.

## 12. PyElly Parsing

Parsing is usually invisible in PyElly operation, which should help to simplify the development of natural language applications. Still, we do sometimes need to look under the hood, either when something goes wrong or when efficiency becomes an issue. So this section will take a deep look at how PyElly parses, which has evolved over its history to become rather complex.

PyElly follows the approach of compiler-compilers like YACC. Compilers are the indispensable programs that translate code written in a high-level programming language like Java or C++ into the low-level machine instructions that a computer can execute directly. In the early days of computing, all compilers were written from scratch; and the crafting of individual compilers was complicated and slow. The results were often unreliable.

To streamline and rationalize compiler development for a proliferation of new languages and new target machines, compiler-compilers were invented. These provided prefabricated and pretested components that could be quickly customized and bolted together to make new compilers. Such standard components typically included a lexical analyzer based on a finite-state automaton and a parser of languages describable by a context-free grammar.

Using a compiler-compiler of course limits the options of programming language designers. They have to work with the constraint of context-free languages; and the individual tokens in that language (variables, constants, and so forth) have to be recognizable by a finite-state automaton. Such restrictions are significant, but being able to develop a reliable compiler in weeks instead of months is so advantageous that almost everyone can live with the tradeoffs.

The LINGOL system of Vaughn Pratt adapted compiler-compiler technology to help build natural language processors. Natural languages are not context-free, but life is more simple if we can parse them as if they were and then take care of context sensitivities through other means like local variables in semantic procedures attached to syntax rules. PyElly follows the LINGOL model and takes it even further.

### 12.1 A Bottom-Up Framework

A parser analyzes an input sentence and builds a description of its structure. As noted earlier, this structure can be represented as a kind of tree, where the root of the tree is a phrase node of the syntactic type `SENT` and the branching of the tree shows how complex structures break down into simpler structures. A tool like PyElly must be able to build such trees incrementally for a sentence, starting either at the bottom with the basic tokens from the sentence or at the top by putting together different possible structures with `SENT` as root and then matching them up with the parts of the sentence.

One can debate whether bottom up or top down is better, but both should produce the same parse tree in the end. We can in fact have it both ways by adopting a basic bottom-



up framework with additional checks to prevent a parse tree phrase node from being generated if it would not also be generated top-down. LINGOL took this approach, and so does PyElly. Going bottom-up here can provide more helpful information when parsing fails, a common occurrence in computational linguistics. It is also more convenient for PyElly cognitive semantics.

The PyElly bottom-up algorithm revolves around a queue that lists the newly created phrase nodes of a parse tree. These still need to be processed to create the phrase nodes at the next higher levels of our tree. Initially, the queue is empty, but we then read the next token in an input sentence and look it up to get some new bottom-level parse tree nodes to prime our queue.

PyElly parsing then runs in a loop, taking the node at the front of its queue and applying its grammar rules to create new nodes to be appended to the back of the queue for further action. This procedure keeps going until the queue finally empties out. At that point, PyElly will then try to read the next token from a sentence to refill the queue and proceed as before. Parsing will stop after every token in sentence has been seen.

There is one circumstance when a new node will not be added to the end of a queue. If there was already a phrase node of the same syntactic type with the same syntactic features built up from the same sentence tokens and if the new node does not have the \*UNIQUE syntactic feature, PyElly will note an ambiguity and will attach the new node as an alternative to the already processed node instead of queueing it separately for further tree building.

This consolidation of new ambiguous nodes serves to reduce the total number of nodes generated for the parsing of a single sentence. Otherwise, PyElly would have to build parallel tree structures for both the old node and the new node without necessarily any benefit. The \*UNIQUE syntactic feature does allow you to override the handling of ambiguities here if you really want to do so.

In any event, PyElly immediately computes the plausibility score of each new phrase as it is generated in parse tree building. Whenever an ambiguity is found, PyElly will find the alternative with the highest plausibility and use it in later processing of a sentence. All the other alternatives, however, will be retained for reporting, for possible backup on a semantic failure, or for adjusting biases to insure that the same rule will not always be taken when there are multiple rules with the same semantic plausibility.

## 12.2 Token Lookup

PyElly token lookup is complicated because it happens in many different ways: external vocabulary tables, FSA pattern rules, entity extraction, and the dictionary rules in a grammar table. These must also interact with macro substitution, inflectional stemming, and morphological analysis; and so we need to understand what is going on here.

PyElly sees a sentence as a sequence of tokens, each a single word or word fragment, a number, a phrase, a complex entity, or a punctuation mark. PyElly parsing goes from

left to right in a sentence, using various language rules to get the extent of the token at the next position. The full procedure is currently as follows:

1. If number rewriting is enabled on, try to rewrite any spelled-out number in the current sentence position as digits plus any ordinal suffix like -ST, -RD, or -TH.
2. Look up the next piece of input text in the external vocabulary table for A; put matches into the PyElly parsing queue as bottom-level phrase nodes if their specified types are consistent with top-down parsing expectations. These expectations are encoded in a binary derivability matrix obtained from current grammar rules.
3. Try also to match up the next piece with the FSA pattern table for A; put matches as phrase nodes into the PyElly parsing queue if their types are consistent with top-down parsing expectations.
4. Try entity extraction at the current position; put any matches as phrase nodes into the PyElly parsing queue if consistent with top-down expectations.
5. If steps 2, 3, or 4 have found pieces of a sentence that are recognizable, find the longest extent and take only phrases of that length as the next token for sentence analysis; discard queued phrase nodes for all shorter extents.
6. Use the default PyElly procedure for extracting the next one-word token at the current sentence position. This may automatically do inflectional stemming and macro substitution.
7. If we already have a token from previous steps that is longer than one word, then stop any further lookup. Proceed instead to the main PyElly parsing loop to work through the phrase nodes already queued up.
8. Look up the next single-word token in the external vocabulary table and in the internal dictionary for a grammar. If found and consistent with top-down expectations, add phrase nodes for the token to the PyElly parsing queue. There will be multiple nodes here if there are multiple vocabulary table entries or grammar table dictionary rules.
9. If we already have any queued phrase nodes to process, then proceed to main PyElly parsing loop.
10. Otherwise, try to break off any morphological endings or beginnings from the current unknown token. If this succeeds, put the reduced token and its affixes back into input text and go back to Step 6 with an empty queue.
11. Check if the next token is known punctuation. If so and punctuation is consistent with top-down expectations, enqueue a phrase node for the punctuation syntactic type `PUNC` and proceed to the main PyElly parsing loop.
12. If all else fails, then create a phrase node of `UNKN` type for the shortest next token and proceed to the main PyElly parsing loop.

## 12.3 Building a Parse Tree

Given a way to put bottom-level phrase nodes into the PyElly parsing queue, we are now ready to build a parse tree from the bottom up. The basic algorithm here is from LINGOL, but the same procedure shows up in other bottom-up parsing systems as well. The next subsection will cover the details of the basic algorithm's main loop, and the two subsections after that will describe PyElly extensions to that algorithm.

---

### 12.3.1 Context-Free Analysis Main Loop

At each step in parsing, we first enqueue the lowest-level phrase nodes for the next piece of an input sentence, with any ambiguities already identified and resolved. Then for each queued phrase node, we go through a process of determining all the ways that the node will fit into a parse tree being built. This is called “ramification” in PyElly source code commentary.

For newly enqueued phrase node, PyElly ramification will go through three steps when the syntactic type of the node is  $X$ :

1. Look for rules of the form  $Z \rightarrow Y \ X$  that have earlier found a  $Y$  and set a goal of an  $X$  in the current position. For each such goal found, create a new phrase node of type  $Z$ , which will be at the same starting position as phrase  $Y$ .
2. Look for rules of the form  $Z \rightarrow X$ . For each such rule, create a new node of type  $Z$  at the same starting position and extent in a sentence as  $X$ .
3. Look for rules of the form  $Z \rightarrow X \ Y$ . For each such rule, set a goal at the next position to look for a  $Y$  to make a  $Z$  at the same starting position as  $X$ .

A new phrase node will be vetoed in steps 1 and 2 if inconsistent with a top-down algorithm. This uses the same binary derivability matrix employed in token lookup.

Each newly created node will be queued up for processing in turn with the three steps above. When all the phrase nodes ending at the current sentence position have been ramified, PyElly parsing advances to the next position.

The main difference between PyElly basic parsing here and similar bottom-up context-free parsing elsewhere is in the handling of ambiguities. Artificial languages generally avoid any ambiguities in their grammar, but natural languages are full of them and so we have to be ready here. In PyElly, the solution is to resolve ambiguities outside of its ramification steps.

PyElly disregards ambiguity when two phrase nodes of different types or features cover the same words in a sentence. For example, the single word **THOUGHT** could be either a noun or the past tense of a verb. This will probably have to be resolved further up in a parse tree, but meanwhile, PyElly will wait for an ambiguity with the same syntactic and features, which will force its hand.

---

### 12.3.2 Special Modifications

Except for ambiguity handling, basic PyElly parsing is quite generic. We can be more efficient here by anticipating how grammar rules for natural language differ from those for context-free artificial languages. The first departure from the core algorithm is the introduction of syntactic features as an extra condition on whether or not a rule is applicable for some aspect of ramification.

On the right side of a rule like  $Z \rightarrow X$  or  $Z \rightarrow X \ Y$ , you can specify what syntactic features must and must not be turned on for a queued phrase node of syntactic type  $X$  to be matched in steps 2 and 3 above and for a queued phrase node of type  $Y$  to satisfy a goal based on a rule  $Z \rightarrow X \ Y$  in step 1. This extra checking has to be added to the basic PyElly parsing algorithm, but it is straightforward to implement.

There is also a special constraint applying to words split into a root and an inflectional ending or suffix (for example, HIT -ING). The parser will set flags in the two resulting nodes so that only step 3 of ramification will be taken for the root part and only step 1 will be taken for each inflection part. A parse tree will therefore grow slower than one might expect here, making for faster parsing.

---

#### 12.3.3 Type 0 Extensions

The introduction of the PyElly . . . syntactic type complicates parsing, but handling the type 0 grammar rules currently allowed by PyElly turns out to require only two localized changes to its core algorithm.

1. Just before processing a new token at the next position of an input sentence, generate a new phrase node for the grammar rule . . . [ . 1 ]  $\rightarrow$  . Enqueue the node and get its ramifications immediately.
2. Just after processing the last token of an input sentence, generate a new phrase node for the grammar rule . . . [ . 2 ]  $\rightarrow$  . Enqueue it and get its ramifications immediately.

Those reading this manual closely will note that the two rules here have syntactic features associated with . . . , which Section 8 said was not allowed. That restriction is still true, and that is because PyElly reserves the syntactic features of . . . to make the type 0 logic handling work properly above.

The difficulty here is that the . . . syntactic type is prone to producing ambiguities. This will be especially bad if the PyElly parser cannot distinguish between a . . . phrase that is empty and one that includes actual pieces of a sentence. So PyElly itself keeps track by using syntactic features here, but keeps that information invisible to users.

The solution will propagate upward the syntactic feature [ . 1 ] that indicates an empty phrase due to case 1 and the feature [ . 2 ] that indicates an empty phrase due to case 2. Though invisible, a grammar will still need to guide this explicitly through setting \*LEFT or \*RIGHT in rules for syntactic feature inheritance when a rule involves . . .

## 12.4 Success and Failure in Parsing

PyElly automatically defines the grammar rule:

```
g:SENT->SENT END
```

---

This rule will never be realized in an actual phrase node, but the basic PyElly parsing algorithm uses this rule to set up a goal for the syntactic type `END` in phase 3 of ramification. After a sentence has been fully parsed, PyElly will look for an `END` goal at the position after the last token extracted from the sentence. If no such goal is found, then we know that parsing has failed; otherwise, we can run the generative semantics for the `SENT` phrase node that generated the `END` goal just found.

There may be more than one `END` goal in the final position, indicating that their respective generating `SENT` phrase nodes were not collapsed as an ambiguity. PyElly can still compare their cognitive semantic plausibility scores, select the most plausible, and run its generative semantic procedure to get the interpretation for a sentence. This is equivalent to making actual phrase nodes based on a `SENT->SENT END` rule, which will trigger PyElly ambiguity handling as just described.

Failure in parsing gives us no generative semantic procedure to run, and our only recourse then is to dump out intermediate results and hope that someone can spot some clue in the fragments. If the failure is due to something discovered in semantic interpretation, though, PyElly can automatically try to recover by backing up in a parse tree to look for an ambiguity and selecting a different alternative at that point.

## 12.5 Parse Data Dumps and Tree Diagrams

PyElly can produce dumps of parsing data, including all the complete or partial parse trees built up for a sentence. In a successful analysis, this helps in verifying that PyElly is running as expected. In a failed analysis, the parsing data will provide clues about what went wrong. For example, you can see where the building of a parse tree had to stop and whether this was due to a missing rule or bad input text.

All this information is written to the standard error stream. Such output originally was an informal debugging aid, but has proved so useful that it is now integral to PyElly operation. The most important part of parse data dumps are the parse trees. These will be presented horizontally, with their highest nodes on the left and with branching laid out vertically. For example, here is a simple 3-level subtree with 4 phrase nodes:

```

sent:0000——ss:8001└noun:8000 @0 [nnnn]
  6 = 3      4 = 2|  1 = 1
                  └verb:0000 @1 [vvvv]
                    2 = -1

```

Each phrase node in a tree display will have the form

```

type:hhhh
  n =  p

```

Where `type` is the name of a syntactic type truncated to 4 characters, `hhhh` is hexadecimal for the associated feature bits (16 are assumed), `n` is phrase sequence number indicating the order in which it was generated, and `p` is the numerical plausibility score computed for the node. The nodes are connected by Unicode drawing characters to show the kind of structural branching.

In the above example, the top-level node here for type `sent` is

```

sent:0000
  6 =  3

```

This node above has no syntactic features turned on; its node sequence ID number is 6, and its plausibility score is +3. Similarly, the node for type `ss` at the next level is

```

ss:8001
  4 =  2

```

The actual sentence tokens for a PyElly will be in brackets on the far right, preceded by its sentence position, which starts from 0. In the example above, the tokens are the “words” `nnnn` and `vvvv` in sentence positions 0 and 1, respectively. Every parse tree branch will end on the far right with a position and token, plus a semantic concept if your grammar includes them.

With analysis of words into components becoming separate tokens, we can get trees like

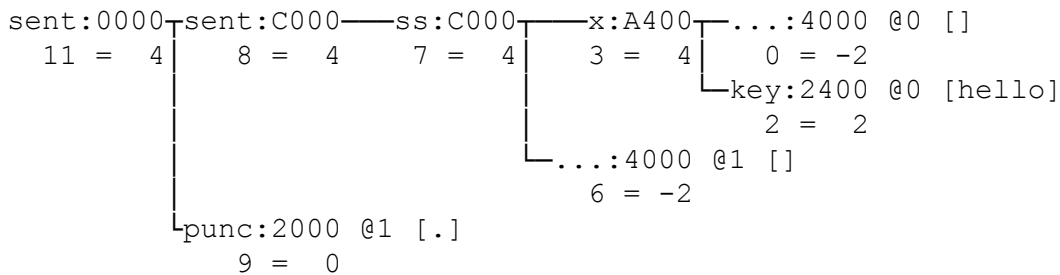
```

sent:0000——ss:0000└ss:0000—unit:0000—unkn:0000 @0 [it]
  11 = 0    10 = 0|  2 = 0    1 = 0    0 = 0
                  └unit:0000└unkn:0000 @1 [live]
                    9 = 0|  4 = 0
                        └sufx:0000 @2 [-s]
                          8 = 0

```

When a grammar includes . . . rules, the display will be slightly more complicated, but still follows the same basic format.

# PyElly User's Manual



The empty phrases number 0 and number 6 have sentence positions 0 and 1, but these are shared by two actual sentence elements HELLO and period (.), as you would expect. Note that the hidden syntactic flags of the . . . type do show up in the displayed tree here; just ignore them.

All the examples here show complete sentences that might be chosen for semantic interpretation. When a PyElly analysis fails, however, you may want to see all the results of parsing, including rejected ambiguities and dead ends. PyElly will accomplish this exploiting the sequence numbers of phrase nodes, which indicate their order of generation.

For a full dump, PyElly finds the node with the highest sequence number not yet shown in any parse tree for the current dump. PyElly then dumps the subtree under that node as done above for the subtree under `sent` and loops back in this fashion until all phrase nodes are accounted for.

In addition to all the trees and subtrees, a PyElly full dump will also show all goals generated in a sentence analysis and all ambiguities found in the process. This information should allow you to reconstruct how PyElly parsed a sentence. Here is an example of a full dump of a short sentence with a trivial grammar:

> **Dogs eat.**

dumping from phrase 5 @0: type=0 [00 00] :0 use=0  
with 4 tokens

```

sent:0000└─noun:0000└─noun:0000 @0 [dog]
      5 = 0      3 = 0      0 = 0
              └─sufx:0000 @1 [-s]
                  2 = 0
              └─verb:0000 @2 [eat]
                  4 = 0

```

dumping from phrase 1 @0: type=4 [00 00] :0 use=0  
with 4 tokens

```

noun:0000 @0 [dog]
  1 = 0

```

0 goals at final position= 4

6 phrases altogether

ambiguities

```

noun 0000: 0 (+0/0) 1 (+0/0)

```

```

raw tokens= [[dog]] [[-s]] [[eat]] [[.]]

```

6 phrases, 5 goals

In this analysis, we have a grammar with the syntactic rule `SENT->NOUN VERB`, the `NOUN DOG` defined in two different dictionary rules, and a dictionary rule for the `VERB EAT`. Parsing fails here for the input “Dogs eat.” because we made no provision for punctuation at the end of a sentence. In the dump, we first see the subtree for the first three tokens of the input sentence, followed by the subtree for the other interpretation of `DOG`. There are no goals at position 4 here. In the listing of ambiguities we have phrases 0 and 1, both identified as a `NOUN` type without any syntactic feature qualifiers.

If a semantic concept is defined for a phrase at a leaf node in a parse tree, a PyElly tree dump will show the concept immediately after the bracketed token at the end of an output line. For example, the augmented printout for the tree above might become

```

sent:0000└─noun:0000└─noun:0000 @0 [dog] 02086723N
      5 = 0      3 = 0      0 = 0
              └─sufx:0000 @1 [-s]
                  2 = 0
              └─verb:0000 @2 [eat] 01170802V
                  4 = 0

```

where 02086723N and 01170802N are WordNet-derived concept names as described in Subsection 10.4.1 above. If no concept is defined for a leaf node, then the tree output will remain the same as before. This is the case for the suffix -S here.



## 12.6 Resource Limits

PyElly is written in Python, a scripting language that can be interpreted on the fly. In this respect, it is closer to the original LINGOL system written in LISP than to any of its predecessors written in Java, C, or FORTRAN. PyElly takes full advantage of Python object-oriented programming and list processing with automatic garbage collection.

Unless you are running on a platform with extremely tight memory, PyElly should be able to handle sentences containing hundreds of tokens with no difficulty. Writing a grammar to describe such huge sentences may take nontrivial effort, however.

The main restrictions in PyElly are the ones on the total number of syntactic types (64) and the total number of different syntactic features for a phrase node (16). These are fixed to allow for preallocation of various arrays used by the PyElly parser in order to get faster operation. You can change those limits in the PyElly Python code, but they should be enough for ordinary applications.

## 13. Developing Language Rules and Troubleshooting

PyElly rewrites input sentences according to the rules that you provide. A natural language application can involve up to eight different `*.*.elly` definition files, however, with some containing hundreds or even thousands of rules. There are plenty of chances to go wrong here; and so a user has to be systematic here and take advantage of all the development aids available.

Even when trying to do something simple, you need to watch out constantly just to avoid falling into bad habits. In general, the best way to use PyElly is to approach a solution by attacking just one sentence at a time. Let PyElly to check everything out for you at each juncture and go no further until everything is satisfactory.

System building will never be a slamdunk, but remember that you are already a natural language expert! Despite enormous advances in hardware and software, an intelligent young child nowadays still knows more about natural language than Siri or Watson. If you can harness some basic analytical skills and programming chops to this innate expertise, then you should do well with PyElly. Just have clear goals and proceed slowly and patiently.

Start with the simplest sentences requiring the fewest rules. Once these can be handled successfully, move on to more complex sentences, adding more rules as needed to describe them. With the PyElly rule framework, you should be able to build on your previous rules without having to change them constantly. This is one big advantage of processing sentences recursively around the syntactic structures of sentences.

When testing out a new sentence, you should not only verify that PyElly is producing the right output, but also inspect its parse tree dump to see that it is consistent with what you expect from your current grammar rules. After any change in rules, you should also check that nothing has broken and that PyElly can still rewrite previously processed sentences. Keep a list of target sentences in a text file to run with `ellyMain.py`; this should be like another PyElly integration test.

### 13.1 Pre-Checks on Language Rule Files

As your language definition files get longer, PyElly can help to verify that each of your separate PyElly language definition files are set up correctly and make sense before you try to bring everything together. This can be done by running the unit tests of the modules to read in definition files and checking for correctness of rules. Running `ellyBase.py` or `ellyMain.py` will also do this to an extent, but you can sort out issues more easily when looking at only one definition file at a time. For example, with application `X`, you can run any or all of the following unit tests from your command line:

```
python grammarTable.py X  
  
python vocabularyTable.py X  
  
python patternTable.py X  
  
python macroTable.py X  
  
python conceptualHierarchy.py X
```

Each command will read in the corresponding `X.*.elly` files, check for errors, and indicate other problems. If a table or a hierarchy can be successfully generated, PyElly will also dump out the entire resulting table. The vocabulary, pattern, and macro table unit tests will also prompt you for specific examples to run against their rules in order to verify correct operation.

PyElly error messages from language definition modules will always be written to `sys.stderr` and will have a first line starting with “\*\* ”. They usually will be followed by a description line starting with “\* ” showing the input text causing the problem. For example, here is an error message for a bad FSA rule for part of speech assignment:

```
** pattern error: bad link  
*   at [ 0 *bbbb* ZED start ]
```

PyElly will continue to process an input file after finding an error so as to catch as many definition problems as possible in one pass.

Once each separate table is defined as far as can be checked in isolation, you can run `ellyBase.py` to load everything together. This will run a cross-table check across on the consistency of your syntactic categories and of your syntactic and semantic features. For application `X`, do this with the command

```
python ellyBase.py X
```

This is also how you would normally test the rewriting of individual test sentences, but the `ellyBase` information from the loading of language rules is a good way to check immediately for omissions or typos in your language rules, which can be hard to track down otherwise. Running `ellyMain.py` will skip over this kind of checking.

## 13.2 A General Application Development Approach

For those wanting explicit details on how to set up PyElly language definition files, here is a reasonable way to build a completely new application *X* step by step.

1. Set up initially empty `X.g.elly`, `X.m.elly`, `X.stl.elly`, and `X.v.elly` files. For the other PyElly language definition files, taking the defaults should be all right.
2. Select representative target sentences for PyElly to rewrite. Five or six should be enough to start with. You will add more as you progress.
3. Write *G*: grammar rules in `X.g.elly` to handle one of your target sentences; leave out the cognitive and generative semantics for now. Just check for correctness of the rules by running the PyElly module `grammarTable.py` with *X* as an argument.
4. Add the words of a target sentence as *D*: internal dictionary rules in `X.g.elly` or as vocabulary entries in `X.v.elly`. Run `grammarTable.py` or `vocabularyTable.py` with *X* as an argument to verify correctness of modified language definition files.
5. Run PyElly module `ellyBase.py` with *X* as an argument to verify that your language definition files can be loaded together. Enter single target sentences as input and inspect the parse data dumps to check for correct analyses. Ignore the generated output for now.
6. Write the generative semantic procedures for your grammar rules and check for correctness by running `grammarTable.py`. If you have problems with a particular semantic procedure, copy its code to a text file and run the PyElly module `generativeProcedure.py` with the name of that text file as an argument.
7. When everything checks out, run `ellyBase.py` with *X* as an argument and verify that PyElly translates a target sentence as you want.
8. When everything is working for a target sentence, add more of them and repeat the above from step 2. Always test your new system against all your old target sentences to make sure that everything is still all right after changes.

## 13.3 Miscellaneous Tips

This subsection is a grab bag of advice about developing nontrivial PyElly grammars and vocabularies based on experience going back to the PARLEZ system, the oldest PyElly ancestor. As with any software development, expect to make mistakes; but try at least to avoid making the old ones over and over.

- PyElly is a simple system of only about sixteen thousand lines of Python code. It is designed to translate certain kinds of strings into other kinds of strings and do nothing more. In other words, it will not by itself help you to replicate Watson or Siri. Go ahead and be ambitious, but be realistic about how much you can accomplish in a short project.
- PyElly analysis revolves around sentences, but remember that you define sentences however you want and need not follow established notions. Try to make your sentences be shorter pieces of text in order to make a PyElly grammar development more manageable.
- The PyElly `ellyMain.py` module is the better choice to rewrite batches of multiple sentences because of its command line options. If you want to work with only one input sentence at a time, run `ellyBase.py`. This is more friendly interactively and also provides more diagnostic information.
- In developing a grammar, keep everything as simple as possible, because you are more likely to run into trouble as the number of syntactic categories and the number of rules increase. In many applications, you can often ignore language details like gender, number, tense, and subject-verb agreement.
- Get the syntax of a target input language right before worrying about the semantics. PyElly automatically supplies you with stubs for both cognitive and generative semantics in grammar rules, which you can replace later with full-fledged procedures.
- Natural language typically has regular and irregular forms. Tackle the regular forms first in your grammar rules and make sure you have a good handle on them before taking on the irregular forms. The latter can often be handled by macro substitutions: for example, change BEAMER to BMW.
- When you build a grammar with semantic procedures, you are programming. Therefore, follow good software engineering practices. Divide a large project into smaller parts that can be finished quickly and individually tested. Test as much as you can as you go along; never wait until all your language rules have been written before testing. The semantic procedures for grammar rules are inherently modular.
- Try to make most of your semantic procedures short, fewer than twenty lines if possible; it will otherwise be hard to verify the correctness of your code visually. If a long procedure is unavoidable, check its code separately with the unit test for the PyElly module `generativeProcedure.py`. Throw in some extra `TRACE` and `SHOW` commands temporarily here to let you monitor what your procedure is doing.
- Be liberal with named generative semantic subprocedures. It can be useful for one to have as few as two or three commands if they will be called more than once. A subprocedure call here could actually take more space than making its commands inline, but clarity and ease of maintenance will trump efficiency here. Spread out

PyElly rewriting to as many separate procedures as you can. Common code used in multiple procedures should always be broken out as named subprocedures.

- It can be helpful to group syntactic types into multiple levels where the semantic procedures at each level will do similar things. Some successive levels might be (1) sentence types, (2) subject and predicate types, (3) noun and verb phrase types, and (4) noun and verb types with inflections. Such levels are also a good way to organize the definition of local variables for communication between different generative semantic procedures.
- Macro substitutions will usually be easier to use than syntax rules plus semantic procedures, but they have to be quite specific about the words that they apply to. Syntax rules are better for patterns that apply to general categories of words.
- Macro substitution rules can be quite dangerous if you are not careful. Watch out for infinite loops of macro substitutions, which can easily arise with \* patterns. Macros can also interact unexpectedly; in particular, make sure that no macro is reversing what another is doing, which may be another cause of infinite loops.
- Try to avoid macros in which the result of substitution is longer than the original substring being replaced. These are sometimes necessary, but can be dangerous; PyElly will warn you if any such macros are detected.
- The ordering of rules in a macro substitution table is important. Rules further up in a list can change the input text that a macro further down the list is looking for.
- PyElly does macro substitution just before the next token is taken from its input buffer and puts results with substitutions back into the input buffer. When inflections and morphological prefixes and suffixes are split off from a word, macros can undo any of this in special cases when necessary.
- For speed, avoid macros for matching literal phrases like “International Monetary Fund.” Unless you need the wild card matching supported by macros, use vocabulary tables instead.
- Macros are powerful, but can slow down processing significantly. This is because all macros have to be checked again after any successful substitution except the null one, and every substitution may involve extensive string copying.
- Vocabulary building should be the last thing you do. You should define at least a few terms to support early testing, but hold off on the bulk of your vocabulary. You first want to know what all your grammar rules look like.
- Syntactic and semantic features can reduce the total number of syntactic rules, but add complexity to your grammar. Use different sets of feature names in different contexts for clarity, but be sure to refer to the right set in any given situation. Once read in by PyElly, all features in rules are stored as anonymous bits and cannot be checked for consistency at run time. Automatic inheritance of syntactic and semantic features through the \*L or \*R mechanisms should be over the same feature sets.

- To dump out the entire saved grammar rule file for application A, run `dumpEllyGrammar.py` with A as an argument. This will also show generative and cognitive semantics, which `ellyBase.py` omits in its diagnostic output.
- Leave the PyElly parse data dumps enabled in PyElly language analysis and learn to read them. This will be the most valuable and easiest to obtain diagnostic information when your language rules are not working as you expect (usually the case). Full tree dumps will show all subtrees generated for ambiguous analyses, but not incorporated into actual PyElly output.
- If you run into a parsing problem with a long sentence, try shortening it to try to isolate the problem. A PyElly parse tree display will be easier to read when working with a shorter sentence.
- When a parse fails, the last token in the listing shown with a parse tree dump will show you where the failure occurred.
- If you are working with English input and have not defined syntax rules for handling the inflectional endings -S, -ED, and -ING, a parse will fail on them. The file `default.p.elly` will define these as `SUFFIX`, but you will then need something in your grammar rules to handle them.
- To follow the execution of a semantic procedure, put a `TRACE` command into it. This will write to the standard error stream whenever it is encountered in a procedure. If a procedure is attached to a phrase node, it will show the syntactic type and starting position of that phrase node and the grammar rule describing the node. If it is in a named subprocedure, PyElly will show the first attached generative procedure calling the subprocedure.
- To see the value of local variables during execution for debugging, use the `SHOW` generative semantic command, which writes to the standard error stream. Remember that both local and global variables will always have string values.
- Punctuation is tricky to handle. Remember that a hyphen will normally be treated as a word break; for example, `GOOD-BYE` currently becomes `GOOD`, `-`, and `BYE`. An underscore or an apostrophe is not a word break, though. Please keep this in mind when writing syntax rules or macro substitutions.
- Macro substitution will not apply across sentence boundaries. To override punctuation otherwise seen as a stop, use the special PyElly stop exception rules described in Subsection 11.1.1. Note that macros can recognize embedded periods and commas, which are non-stopping punctuation.
- Ambiguity is often seen as a problem in language processing, but PyElly embraces it. Deliberate ambiguity can simplify a grammar. For example, the English word `IN` can be either a preposition or a verb particle. Define rules for both usages with plausibility scores and let PyElly figure out which one to apply.

- Try always to give ambiguous alternatives different plausibility scores; otherwise PyElly will switch between them in parsing different sentences, which may not what you want.
- When assigning plausibility scores to rules, try to keep adjustments either mostly positive or mostly negative. Otherwise, they can cancel each other out in unexpected and possibly unfortunate ways because plausibility for a phrase is computed by recursively adding up all the plausibility scores for all its subconstituents.
- Experiment. PyElly offers an abundance of language processing capabilities, and there is often more than one way to do something. Find out what works best for you.
- Fix any language definition problems due to typos first. It is easy to mistype names of syntactic categories or names of semantic or syntactic features. Always check the ellyBase complete listing of grammar symbols in its diagnostic output to verify that there are no unintended ones due to typos.



## 14. PyElly Applications

PyElly by itself is no silver bullet for natural language processing despite its broad range of builtin capabilities. Within its limitations, however, you can produce some quite useful results quite quickly. A good candidate application for PyElly implementation should meet the following conditions:

1. Your input data is UTF-8 Unicode text consisting of either Latin-1 or ASCII characters and divisible into sentences. This need not necessarily be English, but that is where PyElly offers the most builtin capabilities.
2. Your intended output will be translations of fairly short input sentences into arbitrary Unicode in UTF-8 encoding, not necessarily in sentences.
3. No world knowledge is required in the translation of input to output except for what might be expected in a simple dictionary.
4. You can describe the translation process in words and mainly need some support in automating it.
5. Your defined vocabulary is limited enough for you to specify manually with the help of a text editor like vi or emacs, and you can tolerate everything else being treated as the `UNKN` syntactic type.
6. Your computing platform has Python 2.7.\* installed. This will be needed both to develop your language rules and to run your intended application.
7. You are working experimentally and are willing to put up with idiosyncratic non-commercial software.

Familiarity with the Python language will help here, but is not mandatory. You will, however, definitely have to write the code for PyElly cognitive and generative semantics. This is nontrivial work, but is in a highly restricted programming language, which should be straightforward for someone with basic coding experience.

To build a PyElly application `A`, you just need to create its associated language definition files. In the most extreme case, these would include `A.g.elly` (grammar), `A.v.elly` (vocabulary), `A.m.elly` (macro substitutions), `A.p.elly` (syntactic type patterns), `A.ptl.elly` (prefix removal rules), `A.stl.elly` (suffix removal rules), `A.h.elly` (semantic concept hierarchy), and `A.sx.elly` (stop punctuation exceptions). The only mandatory one is `A.g.elly`; the others can be either be empty files or omitted, in which case the respective files for `default` will be loaded instead.

Here are three fairly simple application projects you can try to build as a way of getting to know PyElly:

- A translator from English to pig Latin.
- A bowdlerizer to replace objectionable terms in text with sanitized ones.
- A part of speech tagger for English words just using morphological analysis.

The PyElly distribution includes examples of actual applications serving in part to exercise PyElly fully for integration testing. These fall into two classes: those used for debugging and validation only and those realizing potentially useful functionality. The latter also make up the major part of PyElly integration testing (See Appendix D).

We note the language definition files provided with each example application below. You can look at the rules to get an idea how to write them for your own applications.

Currently, the four basic applications for debugging or validation only are:

**default** (.g,.m,.p,.ptl,.stl,.sx,.v) - not really an application, but a set of language definition files that will be substituted if a particular application does not specify one. These include rules for sophisticated morphological stemming and vocabulary definitions for most of the terms in WordNet 3.0.

**echo** (.g,.m,.p,.stl,.v) - a minimal application that echoes its input as analyzed by PyElly into separate tokens. It will, however, show the effect of inflectional stemming in English on words and entity transformations. You can disable the stemming in your `ellyMain.py` command line.

```
input:  Her faster reaction startled him.
output: her fast -er reaction startle -ed him.
```

**test** (.g,.m,.p,.ptl,.stl,.v) - for basic testing with a vocabulary of short fake words for faster keyboard entry; its grammar defines only simple phrase structures. This was defined as part of integration testing for the first PyElly beta release.

```
input:  nn ve on september 11, 2001.
output: nn ve+on 09/11/2001.
```

**bad** (.g,.h,.m,.p,.v) - deliberately malformed language rules to test PyElly error detection, reporting, and recovery. This is a major part of integration testing, but no grammar or vocabulary table will be generated, and so PyElly will be unable to translate anything here. The files `bad.main.txt` and `bad.main.key` are defined for use with `doTest`, but will always be empty.

```
input:  - -
output: - -
```

The second, more substantial, class of applications are derived from various demonstrations written for PyElly or its predecessors. These have nontrivial examples of PyElly language definitions, illustrate various PyElly capabilities, and provide a basis for

broad integration testing. Most are only prototypes, but you can flesh them out for more operational usage by adding your own vocabulary and grammar rules.

**indexing** (.g,.p,.ptl,.v) - to check removal of purely grammatical words (stopwords), stemming, morphological analysis, and dictionary lookup. Since it obtains roots of content words from arbitrary input text, it could be used to predigest input English text for information searching, statistical data mining, or machine learning systems. This application was written for PyElly.

```
input:  We never had the satisfaction.
output: - - - - satisfy -
```

Note that non-content words are replaced with a hyphen (-).

**texting** (.g,.m,.p,.ptl,.stl,.v) - a test with a big grammar and nontrivial generative semantic procedures. This implements a kind of more or less readable text compression similar to that seen in mobile messaging. It was adapted from a demonstration written for the Jelly predecessor of PyElly.

```
input:  Government is the problem.
output: govt d'prblm.
```

**doctor** (.g,.m,.ptl,.stl,.v) - This has a big grammar with extensive ambiguity handling required. It uses the PyElly . . . syntactic type to emulate Weizenbaum's Doctor program for Rogerian psychoanalysis and was first written to run with the nlf predecessor of PyElly.

```
input:  My mother is always after me.
output: CAN YOU THINK OF A SPECIFIC EXAMPLE.
```

**chinese** (.g,.m,.ptl,.v) - a test of Unicode input and output. It demonstrates some basic Chinese grammar with translations in either traditional [tra] or simplified [sim] characters. This was written to demonstrate PyElly Unicode handling and is still in progress.

```
input:  they sold those three big cars.
output: [sim]他們卖了那三辆大汽车.
output: [tra]他們賣了那三輛大汽車.
```

Only one form of Chinese output will be shown at time in actual operation. You get traditional character output when `ellyMain.py` is run with the flag `-g tra`. The default is simplified output as with `-g sim`. The current integration test is with traditional characters.

**querying** (.g,.m,.ptl,.stl,.v) - heuristically rewrites English queries into SQL commands directed at a relational database of Soviet Cold War aircraft organized into multiple tables. This uses a reworking of language definition files originally written for the PARLEZ and AQF predecessors of PyElly.

```
input:  how high can the foxbat fly?
output: from Ai a,AiPe b
        select ALTD
        where NTNM=foxbat,a.NTNM=b.NTNM
        ;
```

Table and field names are abbreviated: Ai is “aircraft,” AiPe is “aircraft performance,” ALTD is “altitude,” NTNM is “NATO name,” and so forth. The original AQF system aimed to make such names transparent to database users as well as to hide the mechanics of query formation.

**disambig** (.g,.h,.stl,.v) - disambiguation with a PyElly conceptual hierarchy by checking the semantic context of an ambiguous term. This is only a demonstration and not yet a full application like the others listed here. It was written mainly as an integration test focusing on cognitive semantics. Its output is a numerical scoring of semantic relatedness between pairs of possibly ambiguous terms in its input and showing their intersection in a conceptually hierarchy along with the actual WordNet 3.1 concepts assigned to them by PyElly.

```
input:  bass fish.
output: 11 00015568N=animal0n: bass0n/[bass] fish0n/[fish]
```

This uses the PyElly output option to show the plausibility of a translation along with the translation itself, which is right of the second colon (:) in the output line. The plausibility will be left of that colon. The output score here is 11, which is quite high, and the output also includes any concept associated with sentence analysis. The example above shows that the intersection of `bass0n` and `fish0n` is under the WordNet concept `00015568n`, which has the label `animal0n`.

These six example applications show the range of possibilities for simple translation in natural language processing. This is all rather basic, but can still be helpful in producing useful results with text data.

In integration testing, the **test**, **echo**, **indexing**, **texting**, **doctor**, **chinese**, **querying**, and **disambig** applications each have input files `*.main.txt` for `ellyMain` to process. The expected translations here are given by the corresponding files `*.main.key`. To run an actual test with a PyElly application, you can use the bash shell script `doTest` in the PyElly distribution (see Appendix D).

Other PyElly applications being considered in the short term with the current and future versions of PyElly are:

**tagging** - rewrite text with XML tagging. This might be employed as a front end for data mining software.

**name** - extract personal names from input text. This may require some additional entity extraction support in Python.

**translit** - transliterate English words into a non-Latin alphabet or into syllabic or ideographic representation.

As can be seen in the applications already implemented, PyElly can already support a broad range of natural language processing despite its simplicity. With further development of system capabilities, we might also go in other directions like

**madlib** - an implementation of the popular party game. This will require some kind of PyElly template class.

Finally, it would be good if someone could write an application where the input text was in a language other than English. This currently would still have to be encoded in the ASCII and Latin 1 blocks of Unicode, but that would nevertheless allow for French, Spanish, German, Czech, or Hungarian input.

PyElly is a throwback as a natural language system in that it calls for the crafting of large numbers of language rules while current practice favors automatic machine learning. Nevertheless, rule-based approaches can often be helpful, especially if someone else has already worked them out. You might as well exploit prior work to get a leg up in system building, even though you may still plan to rely mostly on machine learning.

PyElly is still a work in progress and probably requires more developing and testing. It is, however, free and compact and can help novices learning about natural language processing or system builders needing to clean up uncontrolled text data. Give it a try. Any criticisms or suggestions here will be welcome, and of course, anyone may freely change the PyElly open-source Python code.

## Appendix A. Python Implementation

This appendix is for Python programmers. You can run PyElly without knowing its underlying implementation, but at some point, you may want to modify PyElly or embed it within some larger information system. The Python source code for PyElly is released under a BSD license, which allows you to change it as needed. You can download it from

`https://github.com/prohippo/pyelly.git`

PyElly was written in Python 2.7.5 under Mac OS X 10.9 and 10.10; it may not run under earlier versions of Python because of changes in the language. To support its external vocabulary tables, PyElly also requires the Berkeley Database (BDB) open-source database manager and the `bsddb3` third-party Python package for accessing the database. Earlier versions of Python had builtin BDB support, but this was dropped because of the difficulty of keeping up compatibility with changes in BDB itself.

Currently, the PyElly v1.0 source code consists of 58 modules, each a text file named with the suffix `.py`. All modules were written to be self-documenting through the standard Python `pydoc` utility. When executed in the directory of PyElly modules, the command

```
pydoc -w x
```

will create an `x.HTML` file describing the Python module `x.py`.

Here is listing of all current PyElly modules grouped by functionality. Some non-Python definition and unit test data files are included in a group when they are integral to the operation or builtin testing of the modules there. These other files are indicated by the dark-shaded rows in the tables below.

## PyElly User's Manual

Inflectional Stemmer (English)	
<code>ellyStemmer.py</code>	base class for inflection stemming
<code>inflectionStemmerEN.py</code>	English inflection stemming
<code>stemLogic.py</code>	class for stemming logic
<code>Stbl.sl</code>	remove -S ending
<code>EDtbl.sl</code>	remove -ED ending
<code>Ttbl.sl</code>	remove -T ending, equivalent to -ED
<code>Ntbl.sl</code>	remove -N ending, a marker of a past participle
<code>INGtbl.sl</code>	remove -ING ending
<code>rest-tbl.sl</code>	restore root as word
<code>spec-tbl.sl</code>	restore special cases
<code>undb-tbl.sl</code>	undouble final consonant of stemming result

Tokenization	
<code>ellyToken.py</code>	class for linguistic tokens in PyElly analysis
<code>ellyBuffer.py</code>	for manipulating text input
<code>ellyBufferEN.py</code>	manipulating text input with English inflection stemming
<code>substitutionBuffer.py</code>	manipulating text input with macro substitutions
<code>macroTable.py</code>	for storing macro substitutions
<code>patternTable.py</code>	extraction and syntactic typing by FSA with pattern matching

## PyElly User's Manual

Parsing	
<code>symbolTable.py</code>	for names of syntactic types, syntactic features, generative semantic subprocedures, global variables
<code>syntaxSpecification.py</code>	syntax specification for PyElly grammar rules
<code>featureSpecification.py</code>	syntactic and semantic features for PyElly grammar rules
<code>grammarTable.py</code>	for grammar rules and internal dictionary entries
<code>grammarRule.py</code>	for representing syntax rules
<code>derivabilityMatrix.py</code>	for establishing derivability of one syntax type from another so that one can make bottom-up parsing do nothing that top-down parsing would not
<code>ellyBits.py</code>	bit-handling for parsing and semantics
<code>parseTreeBase.py</code>	low-level parsing structures and methods
<code>parseTreeBottomUp.py</code>	bottom-up parsing structures and methods
<code>parseTree.py</code>	the core PyElly parsing algorithm
<code>parseTreeWithDisplay.py</code>	parse tree with methods to dump data for diagnostics

Semantics	
<code>generativeDefiner.py</code>	define generative semantic procedure
<code>generativeProcedure.py</code>	generative semantic procedure
<code>cognitiveDefiner.py</code>	define cognitive semantic procedure
<code>cognitiveProcedure.py</code>	cognitive semantic procedure
<code>semanticCommand.py</code>	cognitive and generative semantic operations
<code>conceptualHierarchy.py</code>	concepts for cognitive semantics

Sentences and Punctuation	
<code>ellyCharInputStream.py</code>	single char input stream reading with <code>unread()</code> and reformatting
<code>ellySentenceReader.py</code>	divide text input into sentences
<code>stopExceptions.py</code>	recognize stop punctuation exceptions in text
<code>exoticPunctuation.py</code>	recognize nonstandard punctuation
<code>punctuationRecognizer.py</code>	recognize default sentence punctuation



## PyElly User's Manual

Morphology	
<code>treeLogic.py</code>	binary decision logic base class for affix matching
<code>suffixTreeLogic.py</code>	for handling suffixes
<code>prefixTreeLogic.py</code>	for handling prefixes
<code>morphologyAnalyzer.py</code>	do morphological analysis of tokens

Entity Extraction	
<code>entityExtractor.py</code>	runs Python entity extraction procedures
<code>extractionProcedure.py</code>	some predefined Python entity extraction procedures
<code>simpleTransform.py</code>	basic support for text transformations and handling of spelled out numbers
<code>dateTransform.py</code>	extraction procedure to recognize and normalize dates
<code>timeTransform.py</code>	extraction procedure to recognize and normalize times of day

External Database	
<code>vocabularyTable.py</code>	interface to external vocabulary database
<code>vocabularyElement.py</code>	binary form of external vocabulary record

Top Level	
<code>ellyConfiguration.py</code>	define PyElly parameters for input translation
<code>ellySession.py</code>	save parameters of interactive session
<code>ellyDefinition.py</code>	language rules and vocabulary saving and loading
<code>interpretiveContext.py</code>	handles integration of sentence parsing and interpretation
<code>ellyBase.py</code>	principal module for processing single sentences
<code>ellyMain.py</code>	top-level main module with sentence recognition
<code>dumpEllyGrammar.py</code>	methods to dump out an entire grammar table

Test Support	
<code>parseTest.py</code>	support unit testing of parse tree modules
<code>stemTest.py</code>	test stemming with examples from standard input
<code>procedureTestFrame.py</code>	support unit test of semantic procedures
<code>generativeDefinerTest.txt</code>	to support unit test for building of generative semantic procedures
<code>cognitiveDefinerTest.txt</code>	to support unit test for building of cognitive semantic procedures
<code>suffixTest.txt</code>	to support comprehensive unit test with list of cases to handle
<code>morphologyTest.txt</code>	to support unit test with prefix and suffix tree logic plus inflectional stemming
<code>sentenceTestData.txt</code>	to support unit test of sentence extraction
<code>testProcedure.*.txt</code>	to run with the <code>generativeProcedure.py</code> unit test to verify correct implementation of generative semantic operations

All \*.py and \*.sl files listed above are distributed together in a single directory. The \*.txt files for unit testing will be in a subdirectory `forTesting`.

The first v0.1beta version of the Python code in PyElly was completed in 2013 with some preparatory work done in November and December of 2012. This was an extensive reworking and expansion of the Java code in its Jelly predecessor, making it no longer compatible with Jelly language definition files. PyElly v1.0 graduated from beta status as of December 14, 2014, but other development is still ongoing. See the `README.txt`.

## Appendix B. Historical Background

The natural language tools in PyElly have evolved greatly in the course of being completely rewritten four times in four different languages over the past forty years. Nevertheless, it retains much of the design of the original PDP-11 assembly language implementation of PARLEZ. Writing such low-level code forced simplicity in software architecture, but this actually was advantageous in later ports to different target computing platforms.

The PARLEZ system, for example, had a stripped-down custom programming language for generative semantics because no better alternative was available at the time. That solution, however, has been serviceable for many natural language processing problems and so has been carried along with only a few changes and additions in systems up to and including PyElly. And, yes, arithmetic is still unsupported.

PyElly does depart in major ways even from its immediate predecessor Jelly, however.

- The inflectional stemmer was reorganized to simplify its set of basic operations and to eliminate internal recursive calling. Stemming logic is now set out in text files and can be edited and reloaded at run-time. The number of special cases recognized in English was expanded. The -N and -T inflectional endings were added.
- PyElly morphological analysis was enhanced to allow proper identification of removed prefixes and suffixes as well as just returning stems (lemmas) as done by Jelly. This results in an analytic stemmer, appropriate to a general natural language tool like PyElly. The number of suffix cases recognized in English was expanded to cover many WordNet exceptions and other English irregular forms.
- The syntactic type recognizer was changed to employ an explicit finite-state automaton where transitions are made when an initial part of an input string matches a specified pattern at a state.
- New execution control options were added to generative semantics. Local and global variables were changed to store string values, and list and queue operations were defined for local variables. Deleted buffer text can now be recovered in a local variable. Support for debugging was expanded.
- Semantic concepts were added to cognitive semantics for ambiguity handling. This makes use of a new semantic hierarchy with information derived from WordNet.
- Vocabulary tables were made more scalable by employing the open-source Berkeley Database package to manage persistent external data.
- A new interpretive context class was introduced to coordinate execution of generative semantic procedures.
- Handling of Unicode was improved. UTF-8 is now employed in both PyElly input and output.

- Sentence and punctuation processing is cleaner.
- The PyElly command line interface was reworked to support new initialization and rewriting options.
- Error handling and reporting has been greatly expanded for the definition of language rules. Warnings have also been added for common problems in definitions.
- New unit tests have been attached to major modules.
- New example applications have also been written to serve as integration tests exercising the broad range of PyElly natural language capabilities; older applications from Jelly and earlier systems have been updated to run in PyElly.

Jelly is now retired and will be unavailable. This in part reflects the growing importance of scripting languages like Python versus Java in software development and education.

PyElly is by no means perfect or complete and might be rewritten in yet another programming language as computing practices change. The goal here, however, is less a long-term utopian system than an integrated set of reliable natural language processing tools immediately helpful to students and others building fairly simple natural language systems. Many of these tools may seem old-fashioned to some technologists; but most have had time to mature and prove their usefulness. There is no point in continually having to reinvent or relearn such capabilities.

This PyElly User's Manual rewrites, reorganizes, and greatly extends the earlier one for Jelly, but still retains major parts from the original PARLEZ Non-User's Guide once printed out on an early dot-matrix line printer. Revisions for clarity, accuracy, and completeness are ongoing. Check <https://github.com/prohippo/pyelly.git> for the latest PDF for the manual.

## Appendix C. Berkeley Database

Berkeley Database is an open-source database package available under a GNU license from Oracle Corporation, which acquired the software in 2006 by buying SleepyCat, the company holding its copyright. Berkeley Database is not included in the PyElly distribution. You must obtain it yourself in accordance with all licensing requirements applying to you. There is no restriction on its free use for educational purposes.

This appendix describes how Berkeley Database was downloaded and installed on a MacOS X machine with the Developer tools already set up. The procedure should be more or less applicable to other Unix as well as Linux environments and may be helpful for Windows systems. No testing has actually been done on those platforms, however.

For background on Berkeley Database, see

[http://en.wikipedia.org/wiki/Berkeley\\_DB](http://en.wikipedia.org/wiki/Berkeley_DB)

For software downloads, you can go to the Oracle website

<http://www.oracle.com/technetwork/database/database-technologies/berkeleydb/downloads/index.html>

to get the latest Berkeley Database distribution file. For MacOS X, this would be \*.tar.gz. Unpack it on your system with the shell command

```
tar -xzvf db-*.tar.gz
```

to produce a subdirectory db-\* containing source code that must be compiled and installed using standard Unix development tools like gcc and make. The instructions for doing so on a Unix system can be viewed in a browser by running a MacOS X shell command line to view a Berkeley Database documentation file:

```
open db-*/docs/installation/build_unix.html
```

The procedure should be fairly straightforward for anyone familiar with Unix. An actual MacOS X walkthrough of such compilation and installation can be found on the Web at

<https://code.google.com/p/tonatiuh/wiki/InstallingBerkeleyDBForMac>

To access Berkeley Database from Python, you must next download and install the bsddb3 package from the web. This is available from

<https://pypi.python.org/pypi/bsddb3>

The installation procedure, however, turns out to be quite complicated, and difficult to carry out directly from a command line. The problem is with dependencies where a module A cannot be installed unless module B is first installed. Unfortunately, such dependencies can cascade unpredictably in different environments, so that one fixed set of instructions cannot always guarantee success.

To avoid missteps and all the ensuing frustrations, the easiest solution is use a software package manager that will trace out all module dependencies and formulate a workable installation path automatically. On MacOS X, several package managers are available, but the current favorite is `homebrew`. See this link for general details:

[http://en.wikipedia.org/wiki/Homebrew\\_\(package\\_management\\_software\)](http://en.wikipedia.org/wiki/Homebrew_(package_management_software))

As it turns out, `homebrew` will also handle the installation of Berkeley Database and the upgrading of Python on MacOS X to version 2.7.5 (recommended for `bsddb3`). Here is a web page showing how this all can be done:

<http://stackoverflow.com/questions/16003224/installing-bsddb-package-python>

The Berkeley Database installation for PyElly development was set up using the `brew` and `pip` commands of `homebrew` after an attempted direct installation stalled in a maze of unresolved dependencies.

The `homebrew` package manager is helpful in that it maintains a shared community library of tested installation “formulas” to work with. These resources are specific to MacOS X, however, making `homebrew` inapplicable to Windows or even Linux or other Unix operating systems. If you are running on one of these non-MacOS X platforms, you have to turn to other software package managers; see

[http://en.wikipedia.org/wiki/List\\_of\\_software\\_package\\_management\\_systems](http://en.wikipedia.org/wiki/List_of_software_package_management_systems)

Some of these managers implement parallels to `homebrew` commands, but you will have to check whether they are actually equivalent.

## Appendix D. PyElly System Testing

After any major change to PyElly source code, you should thoroughly validate the resulting system. Check first that every Python module compiles with no errors and then run the current suite of unit and integration tests in the PyElly package.

The following PyElly Python modules have their own unit tests:

cognitiveDefiner	cognitiveProcedure	conceptualHierarchy
conceptualWeighting	dateTransform	derivabilityMatrix
dumpEllyGrammar	ellyBase	ellyBits
ellyBuffer	ellyBufferEN	ellyCharInputStream
ellyDefinition	ellyDefinitionReader	ellySentenceReader
ellyWildcard	entityExtractor	featureSpecification
generativeDefiner	generativeProcedure	grammarRule
grammarTable	inflectionStemmerEN	macroTable
morphologyAnalyzer	parseTree	parseTreeBase
parseTreeBottomUp	parseTreeWithDisplay	patternTable
prefixTreeLogic	punctuationRecognizer	simpleTransform
stemLogic	stopExceptions	substitutionBuffer
suffixTreeLogic	syntaxSpecification	timeTransform
treeLogic	vocabularyTable	

Most of these unit tests are self-contained with predefined input test data, but some also will read sys.stdin to get additional input for testing:

ellyBufferEN	ellyCharInputStream	entityExtractor
inflectionStemmerEN	macroTable	morphologyAnalyzer
patternTable	stemLogic	substitutionBuffer
suffixTreeLogic	vocabularyTable	

Such extra testing is to help you track down a specific problem in one of these modules. You can enter as many additional test examples as you want; just type a <RETURN> by itself to terminate the input loop here. Manually entered test examples will be entirely optional for PyElly validation.

For integration testing, run the following PyElly applications with language definition files included in the `applcn` subdirectory in the PyElly download package:

```
./doTest echo
./doTest test
./doTest indexing
./doTest texting
./doTest doctor
./doTest chinese
./doTest querying
./doTest disambig
```

Compare the actual translated output of each PyElly test application with its expected output as listed in its corresponding `*.main.key` file. These should match up exactly except possibly for extra white space. If you make major changes to PyElly for any reason, you must update all the integration tests affected.

These various applications also are regression tests to verify that a new version of PyElly can still do what it used to. Three of the applications were originally developed to run on predecessors of PyElly, but their language rules have been updated to be consistent with current PyElly conventions and resources.

The integration test applications have been chosen in order to cover a broad range of PyElly processing, although probably not everything yet. New applications will be added to the test suite as appropriate, especially if they are found to break PyElly somehow. This will no doubt happen more often with wider usage of PyElly.

No numbered version of PyElly will be released until it has passed all current unit and integration tests and all its modules have been checked with the `pylint` tool to find any otherwise undiscovered problems. PyElly does not yet conform to all suggested Python coding standards, however, because its design still reflects that of its predecessors.